



# **BMCS3003**

## **Distributed Systems and Parallel Processing**

L03 - Memory Management

Presented by

Assoc Prof Ts Dr Tew Yiqi  
May 2023

# Table of contents

**01**

**Centralised Memory  
Management**

**03**

**Memory Migration**

**02**

**Simple, Shared,  
Distributed Shared  
Memory**

01

# Centralised Memory Management

An architecture in which memory allocations are grouped based on their type, lifetime, or other requirements.



**As computing tasks get larger and larger, may need to enlist more computers to the job.**



Tasks feeder



**As computing tasks get larger and larger, may need to enlist more computers to the job.**

Bigger:

more memory and storage

Faster:

each processor is faster


More tasks:

do many computations simultaneously


# Review of Centralised Memory Management

## Virtual memory

- Extending the size of available memory beyond its physical size of RAM
- Paging versus segmentation
- Internal versus external fragmentation
- Segment placement algorithms
- Page replacement algorithms
  - page faults and thrashing



# Simple Memory Model



In Simple Memory Model access times for all processors are equal.

Requires strict control of degree of multi-programming.

Often does not use virtual memory or caching because of overhead.

# Shared Memory Model


Used for inter-process communication.

Multiple processes share memory locations.


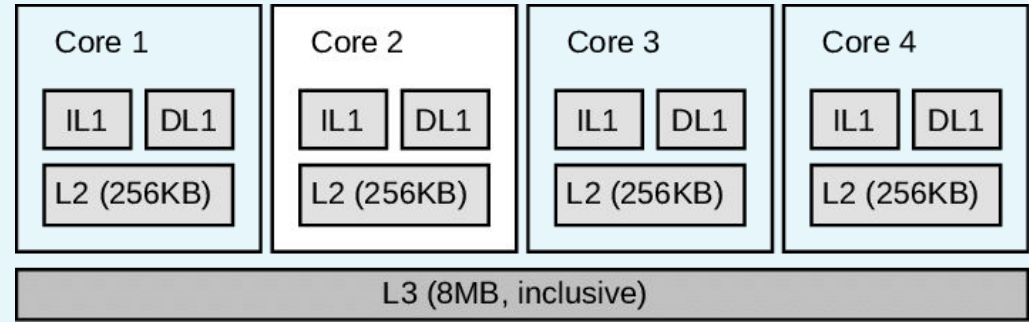

May includes physical RAMs, local cache, and secondary storage.

Memory access takes place via common bus thus a possibility of a **bus contention**.


Example: OpenMP







# Intel Core i7-970 Processor




# Bus Contention

Bus contention occurs when the demand for bus access is excessive.

Bus contention may cause a bottleneck in a shared memory system.

Larger multiprocessor systems (>32 CPUs) cannot use a single bus to interconnect CPUs to memory modules, because bus contention becomes unmanageable.

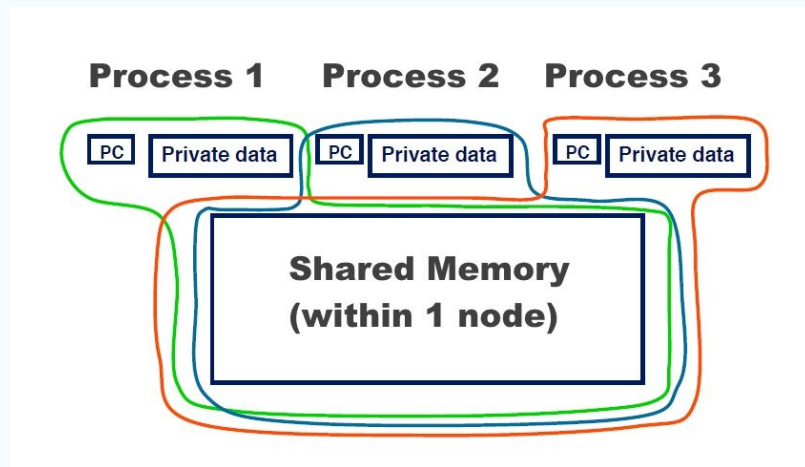


# Shared Memory Performance

Performance is an important issue in a *Shared Memory System*.

Important issues:

- scalability ability to accommodate growth without sacrificing performance
- real time needs
  - overlap of communication and computation
  - prefetching of data
- non local memory references are expensive (up to 10:1 ratio)




More reference: [OpenMP MIMD](#)

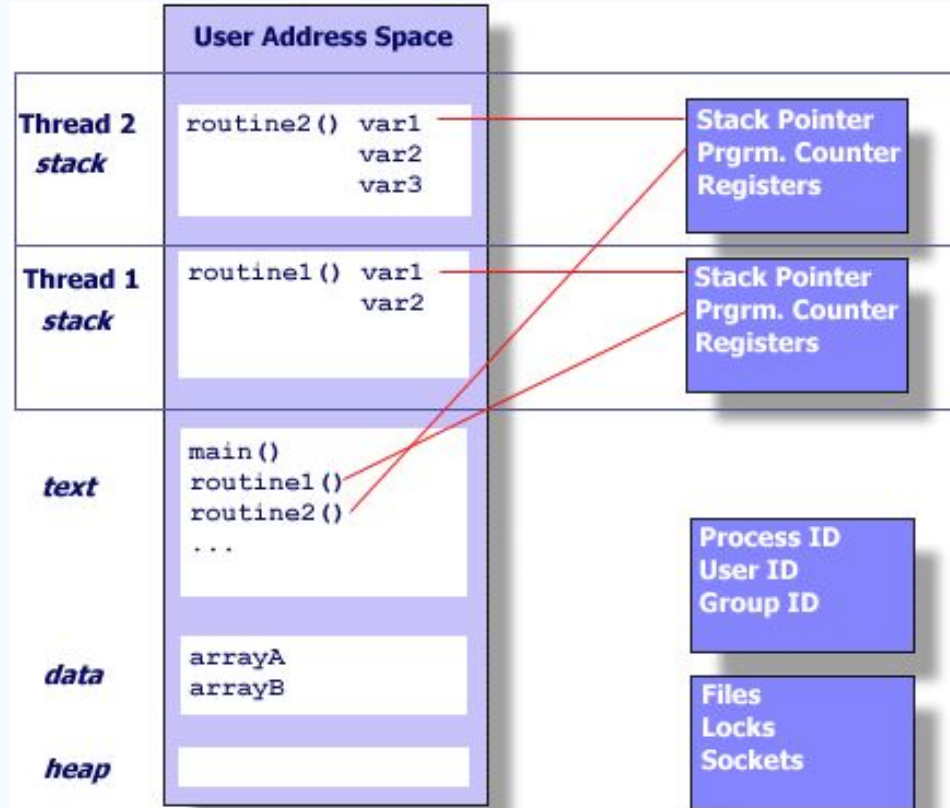
# Shared Memory Performance

Large symmetric multiprocessor systems offered more compute resources to solve large computationally intense problems.

- Threading is the most popular shared memory programming technique
- Advantage:
  - makes it easy for a developer to divide up work, tasks, and data.
- Disadvantage:
  - data races
- E.g., OpenMP



**Threads** are lightweight process and share process state among multiple threads. This greatly reduces the cost of switching contexts




OpenMP is an acronym for **Open Multi-Processing**. OpenMP is a directive-based Application Programming Interface (API) for developing parallel programs on shared memory architectures. The OpenMP standard is maintained by the OpenMP ARB, a corporation whose board of directors includes representatives from many major computer hardware and software vendors.

"The OpenMP ARB (Architecture Review Boards) mission is to standardize directive-based multi-language high-level parallelism that is performant, productive and portable."



Source: [Cornell University](#)



// A pragma defines beginning of a parallel region and specifies iterations of the following for loop should be spread across openmp threads and is the extent of the parallel region


```
#pragma omp parallel for
for (i =0; i < n; i++)
{
    . . . ;
    computations to be completed ;
    . . . ;
}
```

# OpenMP®

Official Documentation  
[OpenMP API](#)

More Lessons:  
[Lawrence Livermore  
National Laboratory](#)

[Jaka's Corner](#)



```
#include <unistd.h>
#include <stdlib.h>
#include <omp.h>
#include <stdio.h>
#define THREADS 4
#define N 16
int main ( ) {
    int i;
    #pragma omp parallel for schedule(static) num_threads(THREADS)
    for (i = 0; i < N; i++) {
        /* wait for i seconds */
        sleep(i);
        printf("Thread %d has completed iteration %d.\n",
            omp_get_thread_num( ), i);
    }
    /* all threads done */
    printf("All done!\n");
    return 0;
}
```


What is this code doing ?

What do the OpenMP semantics specify ?

How might you accomplish this ?

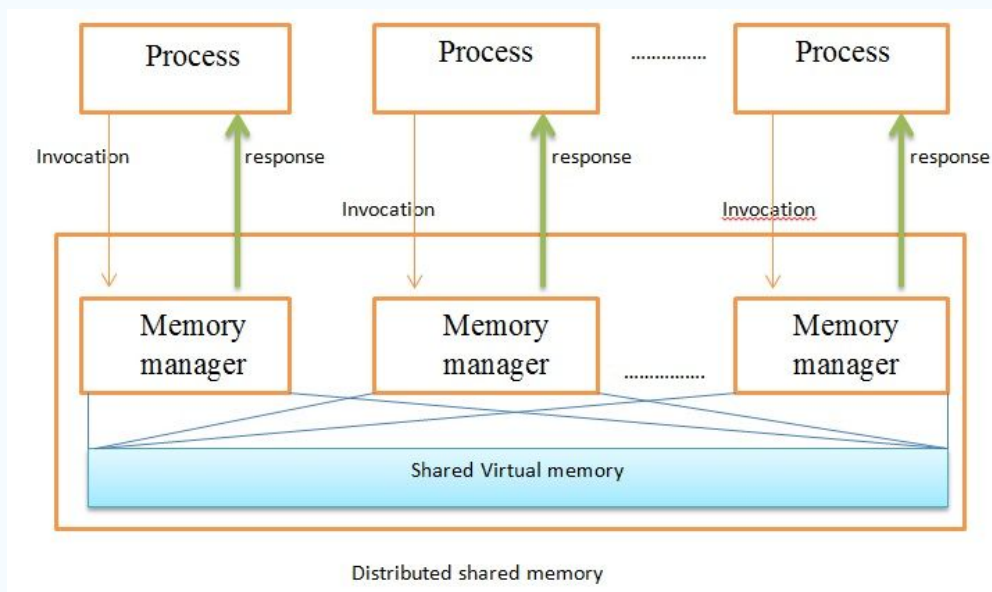
Under the hood:

1. Scheduling
2. Work (in parallel)
3. Reduction
4. Barrier ...





# Distributed Shared Memory (DSM)



Source: [Mehrnazzhian](#)

Concept introduced in 1989.  
DSM presents a logical shared memory for multi-computer systems.

DSM maintains communication and data consistency for applications.

Usually portions of local memory are mapped onto DSM.

DSM maintains a directory service for all data residing in the system.

# Distributed Shared Memory Management

DSM Management involves two main decisions:


- How to distribute shared data ?
- How many readers and writers should be allowed for a shared data segment?



# Distributing Shared Data


## Replication


maintaining multiple  
copies of shared data  
at various locations



## Migration

moving the copy of shared  
data to various locations  
(only one copy of shared  
data is allowed)





# DSM Performance


*Issue to be considered*

## Thrashing

How to avoid thrashing,  
i.e. the situation where  
data constantly travels  
between various locations?

## Data Location

Which is the best original site  
for shared data location?



## Block Size Selection

What is the best size for sharable  
data block?  
(page size, packet size,  
segment size)

## Implementation Location

Where should be DSM  
implemented?  
(hardware or software or both)


# Distributed Memory Programming

Explicitly packaged data and sent it to another system in the cluster.



Message Passing Interface (MPI)

- Explicitly handle the decomposition of the problem across the cluster as well as make sure messages were sent and received in the proper order.



# What is MPI ?



- Message passing is a programming model for coordination processes on these computers
- Each process as its own data
- Data is exchanged by sending and receiving messages
- Programmers use these tools to build multi-computer computation

Official MPI community:

[MPI Forum](#)

Exercises:

[UPM](#)


[LLNL](#)

[Rabernat \(Python\)](#)

[MPI Tutorial \(AWS\)](#)


[Microsoft MPI](#)





C - calls for MPI process 0 to send and receive border data with MPI process 1

```
CALL MPI_ISEND(SBUF, scount, MPI_REAL, 1, stag, MPI_COMM_WORLD, ireq, ierr)
CALL MPI_Irecv(RBUF, rcount, MPI_REAL, 0, rtag, MPI_COMM_WORLD, rreq, ierr)
```



# What is MPI ?



- MPI is a standard library for message passing.
- Ubiquitous in high performance technical computing.
- Nearly every big academic or commercial simulation or data analysis running on multiple nodes uses MPI directly or indirectly.

Official MPI community:  
[MPI Forum](#)


Exercises:  
[UPM](#)

[LLNL](#)

[Rabernat \(Python\)](#)

[MPI Tutorial \(AWS\)](#)

[Microsoft MPI](#)





# Why bother with MPI ?

*Don't we have network libraries?*

- Optimized for performance
- Will take advantage of fastest transport found:
  - Shared memory (Within a computer)
  - Fast cluster interconnects (Infiniband , Myrinet ..)between computers (nodes)
  - TCP/IP if all else failes

Official MPI community:

[MPI Forum](#)

Exercises:


[UPM](#)

[LLNL](#)

[Rabernat \(Python\)](#)

[MPI Tutorial \(AWS\)](#)

[Microsoft MPI](#)



# Why bother with MPI ?

*Don't we have network libraries?*

- Optimized for performance.
- Will take advantage of fastest transport found.
- Enforces certain guarantees:
  - Reliable messages
  - In order arrival of messages

Official MPI community:

[MPI Forum](#)

Exercises:


[UPM](#)

[LLNL](#)

[Rabernat \(Python\)](#)

[MPI Tutorial \(AWS\)](#)

[Microsoft MPI](#)



# Why bother with MPI ?

*Don't we have network libraries?*

- Optimized for performance.
- Will take advantage of fastest transport found.
- Enforces certain guarantees.
- Designed for multi-node technical computing:
  - Completely standard: Available everywhere
  - Has specialized routines for collective operations

Official MPI community:

[MPI Forum](#)

Exercises:


[UPM](#)

[LLNL](#)

[Rabernat \(Python\)](#)

[MPI Tutorial \(AWS\)](#)

[Microsoft MPI](#)



# MPI basic functions

## Set up / teardown

MPI\_INIT  
MPI\_FINALIZE

## Who am I?

MPI\_COMM\_SIZE  
MPI\_COMM\_RANK

## Message Passing

MPI\_SEND  
MPI\_RECV

More Examples:

### PDC-support

By Jarno Rantaharju et al.

### Universiti of Minnesota

By David Porter &  
Mark Nelson

### Princeton Plasma Physics Lab

By Stephane Ethier

# Question


**So is MPI the way  
to do all distributed  
computing ?**



# A Simple MPI Program

```
#include "mpi.h"
#include <stdio.h>

Int main (int argc, char *argv[]){
    Int rank, size ;
    MPI_Init(NULL NULL);
    MPI_Comm_rank(MPI_COMM_WORLD , &rank);
    MPI_Comm_size(MPI_COMM_WORLD , &size);
    printf("Hello! I am %d of %d\n", rank, size );
    MPI_Finalize();
    return 0;
}
```




- **MPI\_COMM\_WORLD** indicates the set of all processes
- All MPI functions return error code
- Update values by passing pointers as arguments

# A Simple MPI Program

```
#include "mpi.h"
#include <stdio.h>
Int main (int argc, char *argv[]){
    Int rank, size ;
    MPI_Init(NULL NULL);
    MPI_Comm_rank(MPI_COMM_WORLD , &rank);
    MPI_Comm_size(MPI_COMM_WORLD , &size);
    printf("Hello! I am %d of %d\n", rank, size );
    MPI_Finalize();
    return 0;
}
```

**- Compile and Run with cmd:**  
`mpi exec np 2 ./hello`



```
Hello! I am 0 of 2
Hello! I am 1 of 2
```

# What just happened ?

- `mpiexec` launched 4 processes on your computer
- Each process ran the program `hostname`
- Each ran independently
- Can do the same thing with 8, 27, or 93 processes
- For performance reasons, usually use as many processes as there are processors





## ... but Painful !

### OpenMP

- Add pragmas to existing program
- Compiler + runtime system arrange for parallel execution
- Rely on shared memory for communication


### MPI

- Must rewrite program to describe how single process should operate on its data and communicate with other processes.
- Explicit data movement: programmer must say exactly what data goes where and when.
- Advantage: Can operate on systems that don't have shared memory.

# MPI - Process Identification

## When running with P processes:

- Size : P  
Total number of processes
- Rank: Number between 0 and P 1  
Identity of individual process
- Library Functions  
MPI\_Comm\_size  
MPI\_Comm\_rank



# Comparing Frameworks



## Multiple Processes

### MPI

Fixed processes number created  
All execute code starting with main  
Isolated address spaces



### FORK

Processes created during program execution  
Replicate address space upon creation,  
but then isolated



### pthread\_create

Threads created during program  
execution  
Shared address space

## Multiple Threads



### OpenMP


Set of threads created at beginning of program  
Recruited to execute tasks spawned by  
`#pragma omp parallel`  
Shared address space

# Synchronous Sending and Receiving

**send():** call returns when sender receives acknowledgement that message data resides in address space of receiver.

**recv():** call returns when data from received message is copied into address space of receiver and acknowledgement sent back to sender

Potential for deadlock if all processes attempt to send and then receive



# Synchronous Sending and Receiving

Sender:

Call **SEND**(foo)

**Copy data** from buffer 'foo' in sender's address  
space into network buffer

Send **message**

Receive **ack**  
**SEND()** returns

Receiver:


Call **RECV**(bar)

Receive **message**  
**Copy data** into buffer 'bar' in receiver's  
address space

Send **ack**  
**RECV()** returns



# Asynchronous Sending and Receiving



## **send(): call returns immediately**

- Buffer provided to send() cannot be modified by calling thread since message processing occurs concurrently with process execution
- Calling thread can perform other work while waiting for message to be sent.

## **recv(): posts intent to receive in the future, returns Immediately**

- Use checksend(), checkrecv() to determine actual status of send/receipt
- Calling thread can perform other work while waiting for message to be received

# Asynchronous Sending and Receiving

Sender:

Call `SEND(foo)`  
`SEND` returns handle `h1`

Copy data from 'foo' into network buffer

Send message

Call `CHECKSEND(h1)` // if message sent, now safe for thread to modify 'foo'

Receiver:

Call `RECV(bar)`  
`RECV(bar)` returns handle `h2`

Receive message

Messaging library copies data into 'bar'

Call `CHECKRECV(h2)`  
// if received, now safe for thread  
// to access 'bar'

RED TEXT = executes concurrently with application thread

# MPI Send / Receive Operation



## Synchronous

MPI\_Send

MPI\_Recv



## Asynchronous

MPI\_Isend

MPI\_Irecv

MPI\_Wait



# Memory Migration

Memory Migration requires two fundamental decisions:

- When in the migration process will we migrate memory?
- How much memory needs to be migrated?



# Vector Add Example

Example

## Steps to parallelize Vector Add with GPU

1. Make the variables accessible to both CPU and GPU
2. Launch configuration of the kernel function
3. Parallelize the kernel function to run on GPU

# More Examples

Example

## Matrix Transpose

<https://www.hpc.cineca.it/content/exercise-15>

## Matrix Multiplication

<https://www.hpc.cineca.it/content/exercise-16>

## 2D Laplace Equation

[https://www.mcs.anl.gov/research/projects/mpi/tutorial/mpie  
xmpl/src/jacobi/C/main.html](https://www.mcs.anl.gov/research/projects/mpi/tutorial/mpie<br/>xmpl/src/jacobi/C/main.html)

# Summary of Chapter



- **Shared address space**


- > Communication is unstructured, implicit in loads and stores.
- > Natural way of programming, but can shoot yourself in the foot easily.
- > Program might be correct, but not perform well.

- **Message passing**

- > Structure all communication as messages.
- > Often harder to get first correct program than shared address space.
- > Structure often helpful in getting to first correct, scalable program.

- **Data parallel**

- > Structure computation as a big “map” over a collection.
- > Assumes a shared address space from which to load inputs/store results, but model severely limits communication between iterations of the map (goal: preserve independent processing of iterations).
- > Modern embodiments encourage, but don’t enforce, this structure.



# Thank you

CREDITS: This presentation template was created  
by **Slidesgo**, and includes icons by **Flaticon**, and  
infographics & images by **Freepik**