



# **BMCS3003**

## **Distributed Systems and Parallel Computing**

L04 - Concurrency Control (Part 1)

Presented by

Assoc Prof Ts Dr Tew Yiqi  
May 2023

# Table of contents

**01**

**Mutual Exclusion &  
Critical Region**

**02**

**Locks**

**03**

**Semaphores**

01

# Mutual Exclusion & Critical Region

A program object that blocks multiple users from accessing the same shared variable data at the same time.

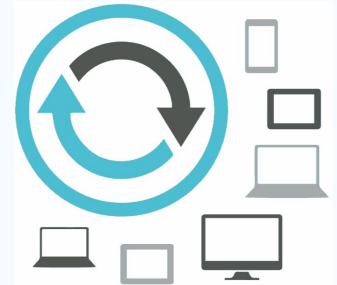


# Synchronization

**Parallel processing (multi-processing)** – 2 or more processors operate in unison.

- 2 or more CPUs are executing instructions simultaneously.
- Each CPU can have a process in RUNNING state at same time.
- Processor Manager has to coordinate activity of each processor and synchronize interaction among CPUs.

**Synchronization** is critical to a system's success because many things can go wrong in a multiprocessing system.



## Example - Application of concurrent processing

Compute  $A = 3 * B * C + 4 / (D+E) ** (F-G)$

COBEGIN

T1 = 3 \* B

T2 = D+E

T3 = F-G

COEND

COBEGIN

T4 = T1 \* C

T5 = T2 \*\* T3

COEND

A = T4+4/T5

CPU 1

CPU 2

CPU 3

CPU 1

CPU 2



## Example - Application of concurrent processing

DO i = 1, 3

A(i) = B(i) + C(i)

ENDDO

Easy ?



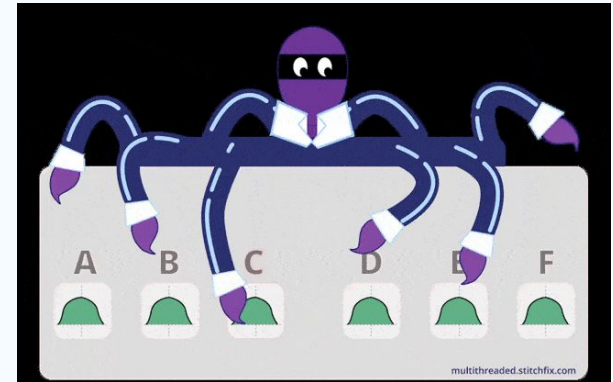
Processor 1 perform :  $A(1) = B(1) + C(1)$

Processor 2 perform :  $A(2) = B(2) + C(2)$

Processor 3 perform :  $A(3) = B(3) + C(3)$

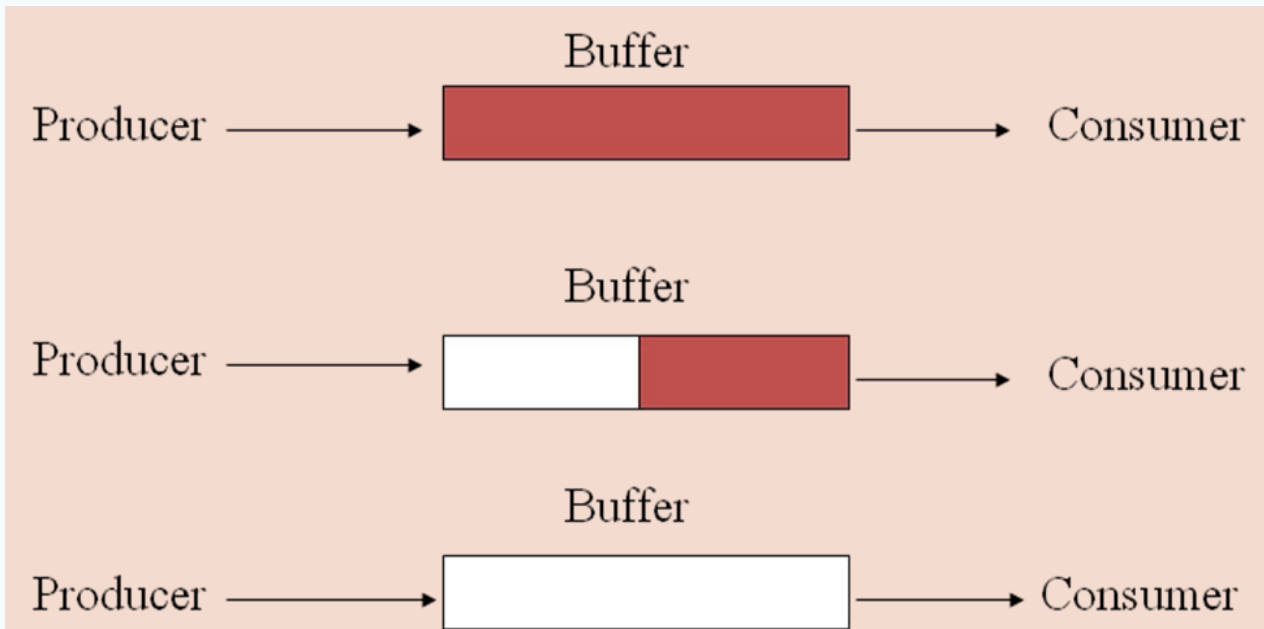
# Classical Synchronisation Problems

- Common element in all synchronization schemes is to allow a process to finish work on a **critical region** of program before other processes have access to it.
- Applicable to both multiprocessors and to 2 or more processes in a single-processor (time-shared) processing system.
- Called a critical region because its execution must be handled as **ONE UNIT**.



# Producers and Customers

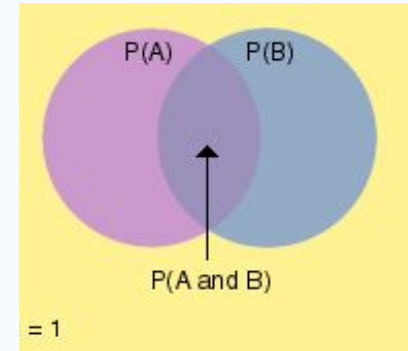
## One Process Produces Some Data That Another Process Consumes Later





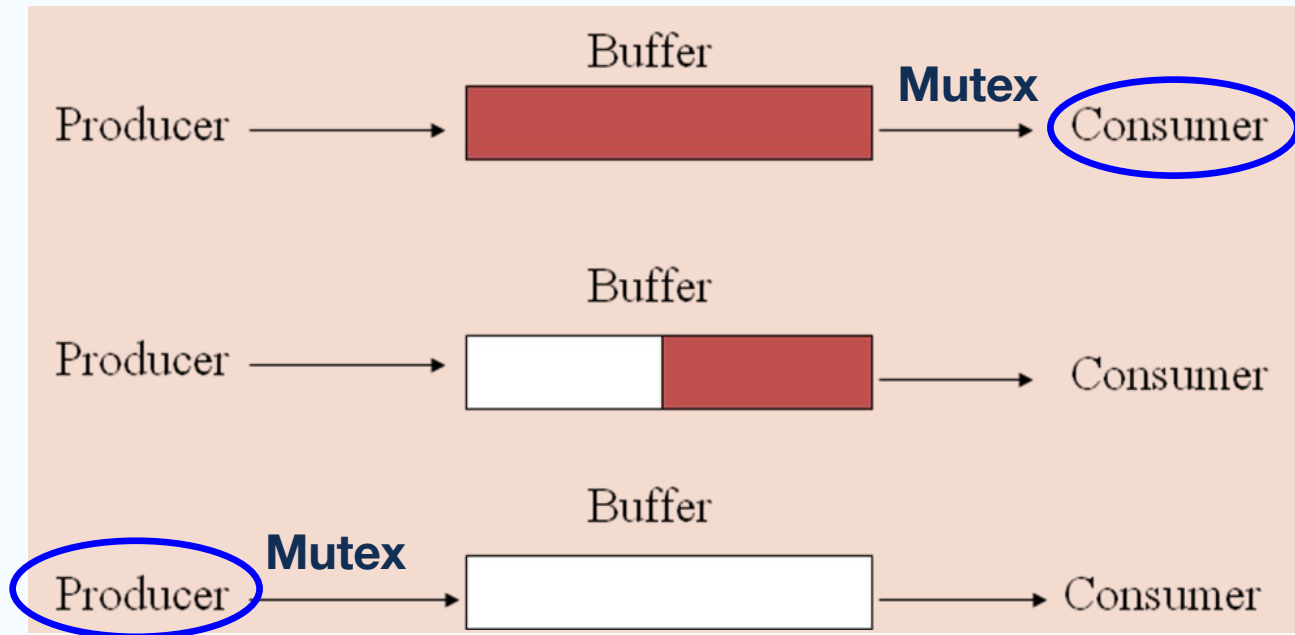
# Mutual Exclusion (Mutex)

- Because buffer holds finite amount of data, synchronization process must delay producer from generating more data when buffer is full.
- Delay consumer from retrieving data when buffer is empty.
- This task can be implemented by 3 semaphores (signals):
  - Indicate number of full positions in buffer.
  - Indicate number of empty positions in buffer.
  - Mutex, will ensure mutual exclusion between processes (A **mutex** is a program object that is created so that multiple program thread can take turns sharing the same resource)



# Producers and Customers

One Process Produces Some Data That Another Process Consumes Later



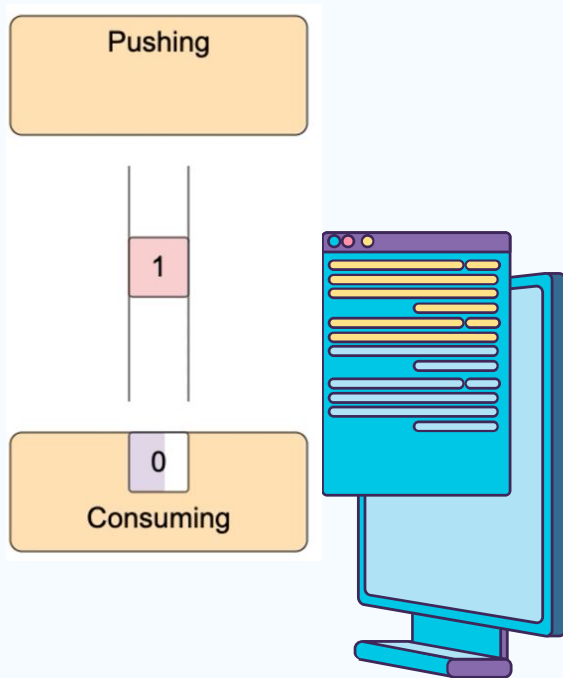
# Producers and Customers

- **Concurrent access** to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Suppose that we wanted to provide a solution to the **consumer-producer problem** that fills **all** the buffers. We can do so by having an integer **count** that keeps track of the number of full buffers. **Initially, count is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.**



# Producers

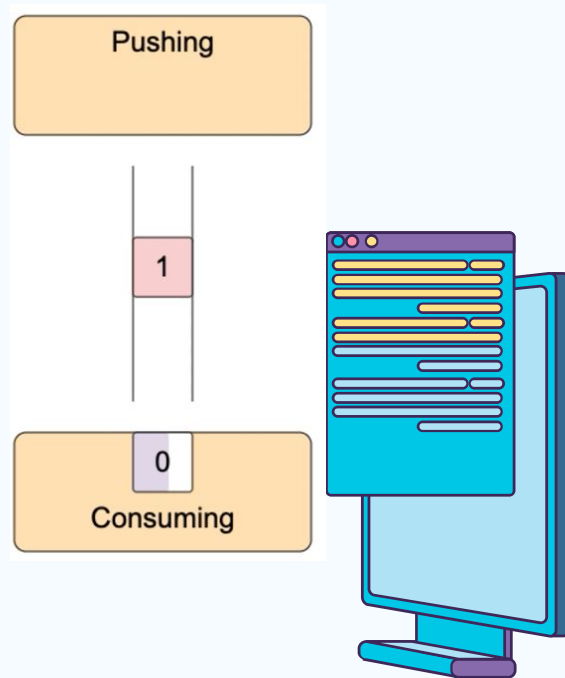
```
while (true)
{
    -- Produce an item and put in nextProduced
    while (count == BUFFER_SIZE)
        ; // do nothing
    buffer [in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    count++;
}
```



# Customers

```
while (true)
{
    while (count == 0)
        ; // do nothing

    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;
    -- Consume the item in nextConsumed
}
```



count

0

in

A

0

1

2

3

out

**Producer**

```
while (true)    {  
    /* produce an item and put in nextProduced  
    while (count == BUFFER_SIZE)  
    ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

**count**

1

in



**Buffer**

A

B

0

1

2

3

out



**Producer**

```
while (true)    {  
    /* produce an item and put in nextProduced  
    while (count == BUFFER_SIZE)  
    ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

**count**

2

in

**Buffer**

A

B

C

0

1

2

3

out

**Producer**

```
while (true)    {  
    /* produce an item and put in nextProduced  
    while (count == BUFFER_SIZE)  
    ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```



**count**

3

in

**Buffer**

A

B

C

0

1

2

3

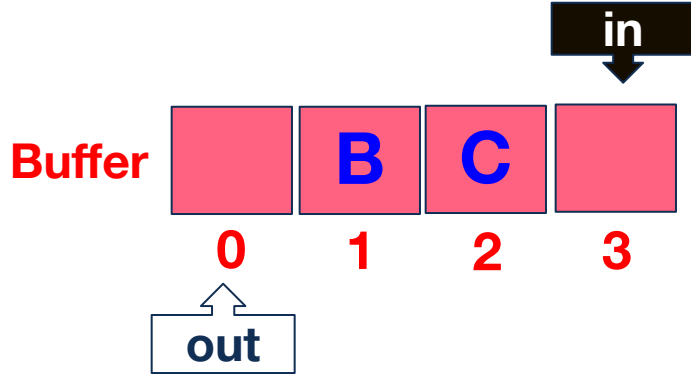
out

**Producer**

```
while (true)    {  
    /* produce an item and put in nextProduced  
    while (count == BUFFER_SIZE)  
    ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

**count**

3

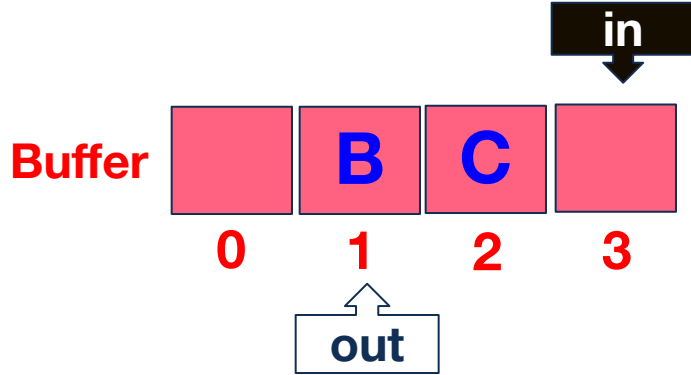


**Consumer**

```
while (true)    {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in nextConsumed  
    }  
}
```

**count**

2



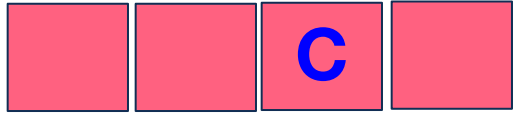
**Consumer**

```
while (true)    {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in nextConsumed  
    }  
}
```

**count**

1

**Buffer**



0

1

2

3

out

**Consumer**

```
while (true)    {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in nextConsumed  
    }  
}
```

**count**

0

**in**



**Buffer**



0

1

2

3

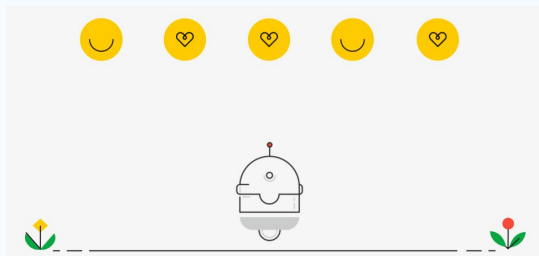
**out**



**Consumer**

```
while (true)    {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed =  buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in nextConsumed  
    }  
}
```

# a. Producer-Consumer Race Condition



➤ `count++` is actually implemented as  
`register1 = count`  
`register1 = register1 + 1`  
`count = register1`

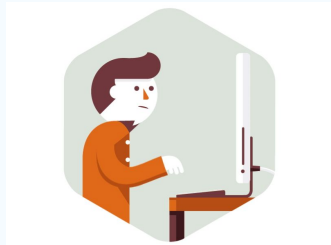
➤ `count--` is actually implemented as  
`register2 = count`  
`register2 = register2 - 1`  
`count = register2`

Consider this execution interleaving with “count = 5” initially:

- t0: producer execute `register1 = count`      {register1 = 5}
- t1: producer execute `register1 = register1 + 1`      {register1 = 6}
- t2: consumer execute `register2 = count`      {register2 = 5}
- t3: consumer execute `register2 = register2 - 1`      {register2 = 4}
- t4: producer execute `count = register1`      {count = 6 }
- t5: consumer execute `count = register2`      {count = 4}

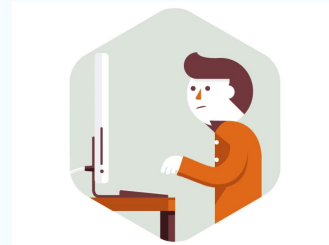
This is a  
problem!  
Count should  
be always 5!

# Illustration



Count

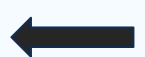
5



**Producer**

**Consumer**

```
register1 = count
register1 = register1 + 1
count = register1
```



**t0**

```
register2 = count
register2 = register2 - 1
count = register2
```

# Illustration

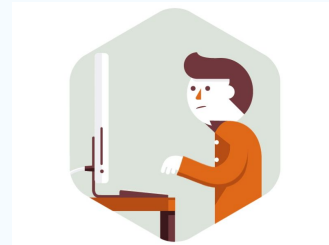


**Producer**

```
register1 = count
register1 = register1 + 1 ← t1
count     = register1
```

**Count**

**5**



**Consumer**

```
register2 = count
register2 = register2 - 1
count     = register2
```



# Illustration

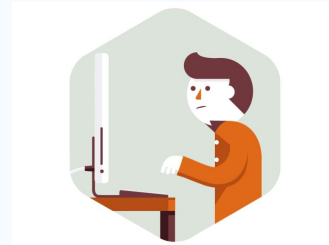


**Producer**

```
register1 = count
register1 = register1 + 1
count     = register1
```

**Count**

**5**



**Consumer**

**t2** ➡

```
register2 = count
register2 = register2 - 1
count     = register2
```

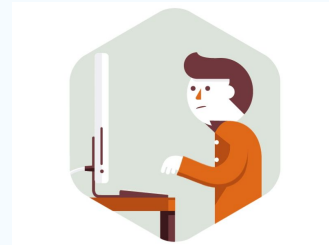
# Illustration



**Producer**

**Count**

**5**



**Consumer**

register1 = count  
 register1 = register1 + 1  
 count = register1

**t3** ➡

register2 = count  
 register2 = register2 - 1  
 count = register2

# Illustration

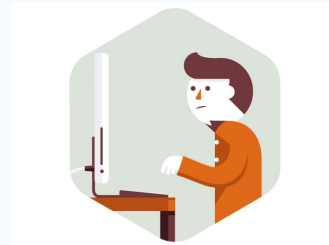


**Producer**

```
register1 = count
register1 = register1 + 1
count = register1
```

**Count**

6



**Consumer**

```
register2 = count
register2 = register2 - 1
count = register2
```

← **t4**

# Illustration

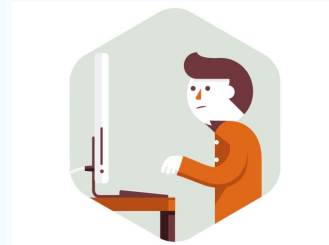


**Producer**

```
register1 = count
register1 = register1 + 1
count = register1
```

**Count**

4



**Consumer**

```
register2 = count
register2 = register2 - 1
count = register2
```

**t5** ➡

# Elaboration

- If the producer and consumer run concurrently the value of the variable counter may be 4 or 6 (the value of counter is 6 when we reversed the order of the statements at t4 and t5)
- Incorrect state occurs because we allowed both processes to manipulate the variable **counter** concurrently.
- Several processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which access take place is called a **Race Condition**

# Elaboration

- In fact, the producer program has to finish its execution then the consumer program to print or display the result, so the counter value suppose is 5. Which is generated correctly if producer and consumer execute separately.



count

5

Producer

```
while (true)    {  
    /* produce an item and put in  
    nextProduced  
    while (count == BUFFER_SIZE)  
    ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

in



Buffer

A

B

C

D

E

F

0

1

2

3

4

5

6

out



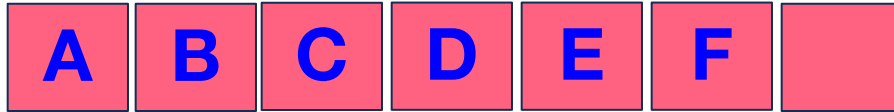
count

6

**Producer**

```
while (true)    {  
    /* produce an item and put in  
    nextProduced  
    while (count == BUFFER_SIZE)  
    ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

Buffer



0

1

2

3

4

5

6

out

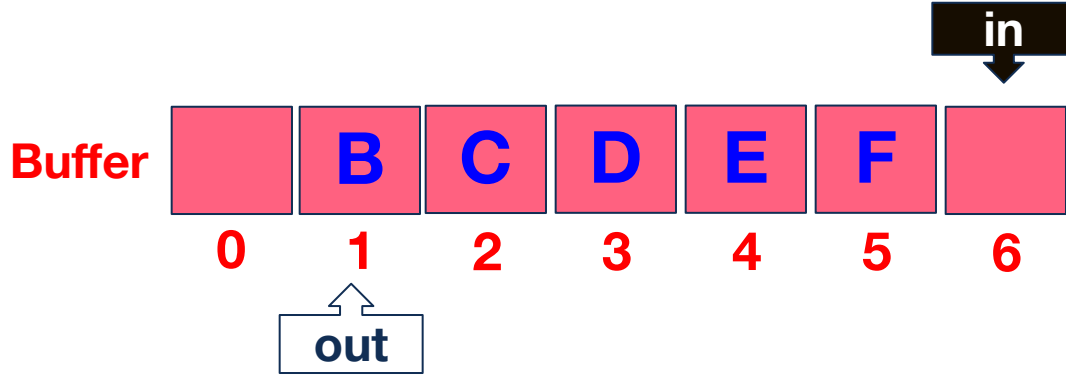


**count**

5

**Consumer**

```
while (true)    {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in  
       nextConsumed  
    */  
}
```



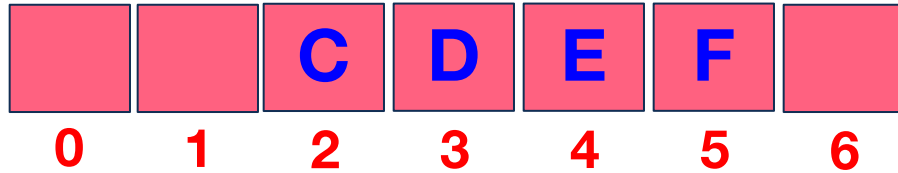
**count**

4

**Consumer**

```
while (true)    {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in  
    nextConsumed  
    }  
}
```

**Buffer**



0

1

2

3

4

5

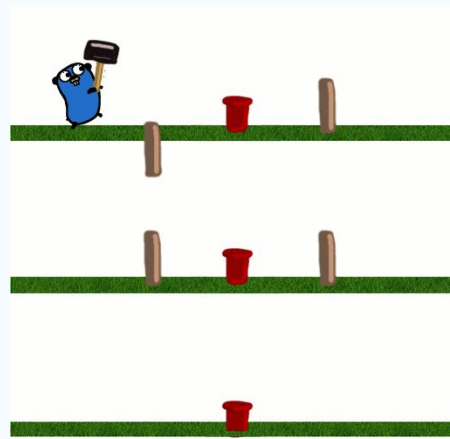
6

out


## b. Print Spooler Race Condition

- A process wants to print a file, it enter the file name in **spooler** directory.
- Another process, the printer **daemon\***, periodically checks to see if there are any files to be printed, then prints them and removes their names from the directory.
- Imagine that our spooler directory has very large number of slots, numbered 0,1,2,3,..., each one capable of holding a file name.
- There are two shared variables **out** and **in**.
  - **out** which points to the next file to be printed.
  - **in** is point to next free slot in directory.
- Process A and Process B decide they want to queue a file for printing.

\* A **daemon** is a program that waits for another program to ask it to do something.



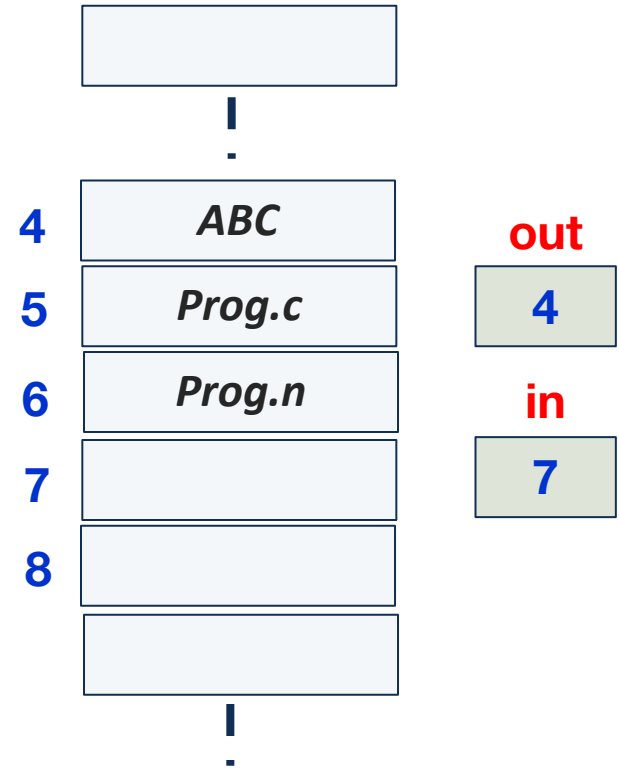
## Process A

read **in**   
store into A\_NextFreeSlot  
add **file a** into the buffer  
**in** = A\_NextFreeSlot + 1

## Process B

read **in**  
store into B\_NextFreeSlot  
add **file b** into the buffer  
**in** = B\_NextFreeSlot + 1

## Spooler Directory/ buffer



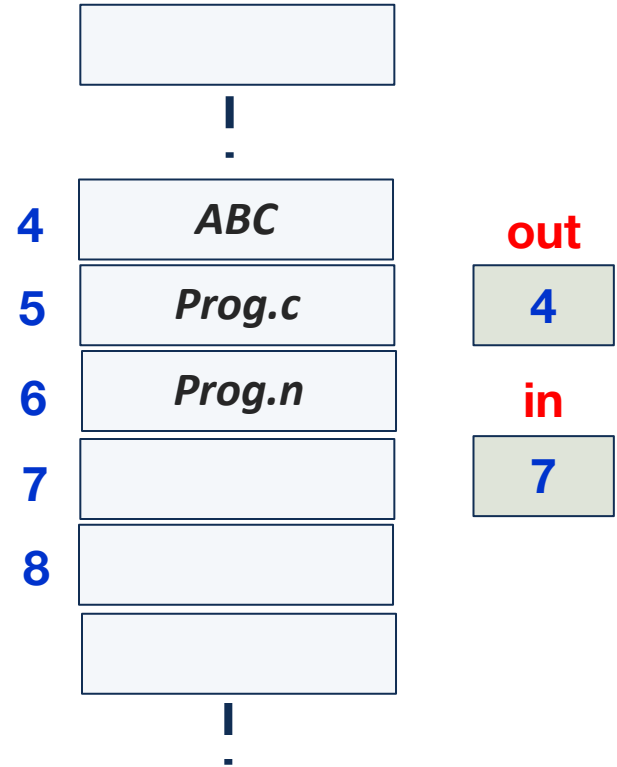
## Process A

read **in**  
store into A\_NextFreeSlot<sup>7</sup> ←  
add **file a** into the buffer  
**in** = A\_NextFreeSlot + 1

## Process B

read **in**  
store into B\_NextFreeSlot  
add **file b** into the buffer  
**in** = B\_NextFreeSlot + 1

## Spooler Directory/ buffer



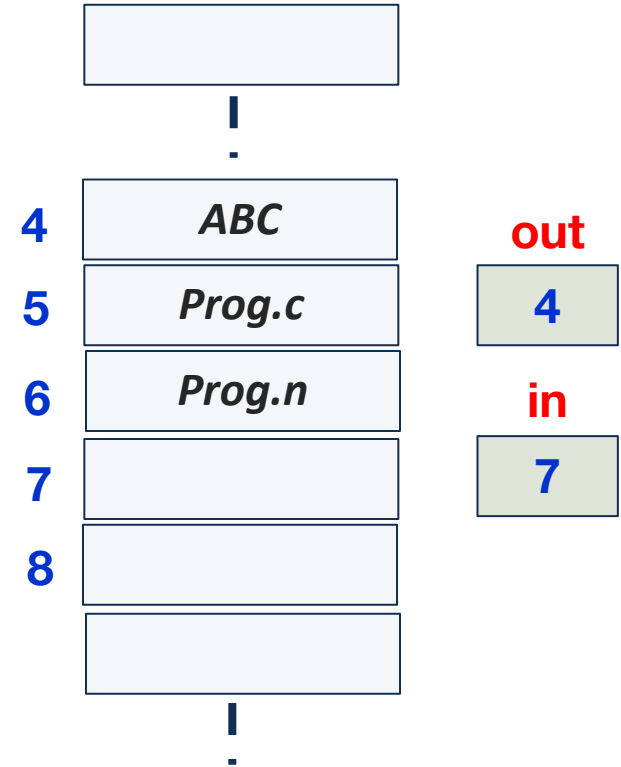
## Process A

read **in**  
store into A\_NextFreeSlot<sup>7</sup>  
add **file a** into the buffer  
**in** = A\_NextFreeSlot + 1

## Process B

read **in**  
store into B\_NextFreeSlot  
add **file b** into the buffer  
**in** = B\_NextFreeSlot + 1

## Spooler Directory/ buffer



## Process A

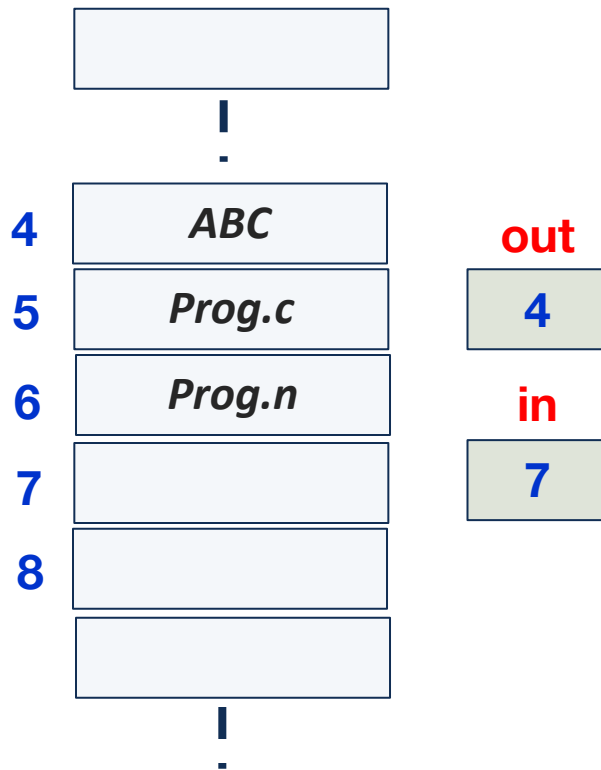
read **in**  
store into A\_NextFreeSlot<sup>7</sup>  
add **file a** into the buffer  
**in** = A\_NextFreeSlot + 1

## Process B

read **in**  
store into B\_NextFreeSlot  
add **file b** into the buffer  
**in** = B\_NextFreeSlot + 1



## Spooler Directory/ buffer



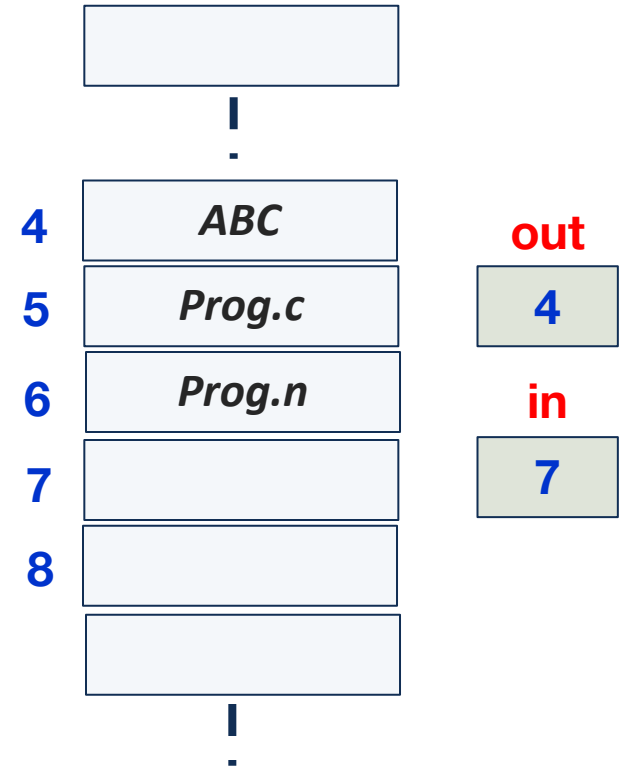
## Process A

read **in**  
store into A\_NextFreeSlot<sup>7</sup>  
add **file a** into the buffer  
**in** = A\_NextFreeSlot + 1

## Process B

read **in**  
store into B\_NextFreeSlot<sup>7</sup> ←  
add **file b** into the buffer  
**in** = B\_NextFreeSlot + 1

## Spooler Directory/ buffer





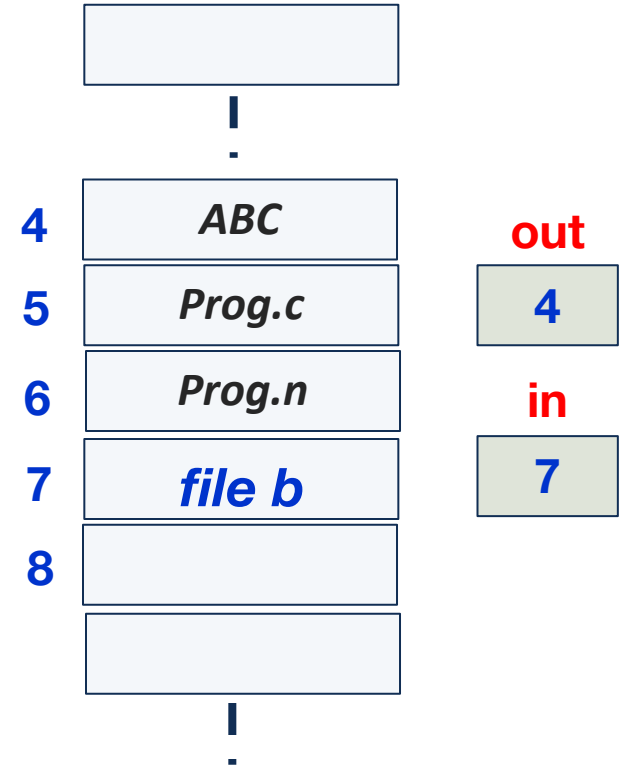
## Process A

read **in**  
store into A\_NextFreeSlot<sup>7</sup>  
add **file a** into the buffer  
**in** = A\_NextFreeSlot + 1

## Process B

read **in**  
store into B\_NextFreeSlot<sup>7</sup>  
add **file b** into the buffer ←  
**in** = B\_NextFreeSlot + 1

## Spooler Directory/ buffer



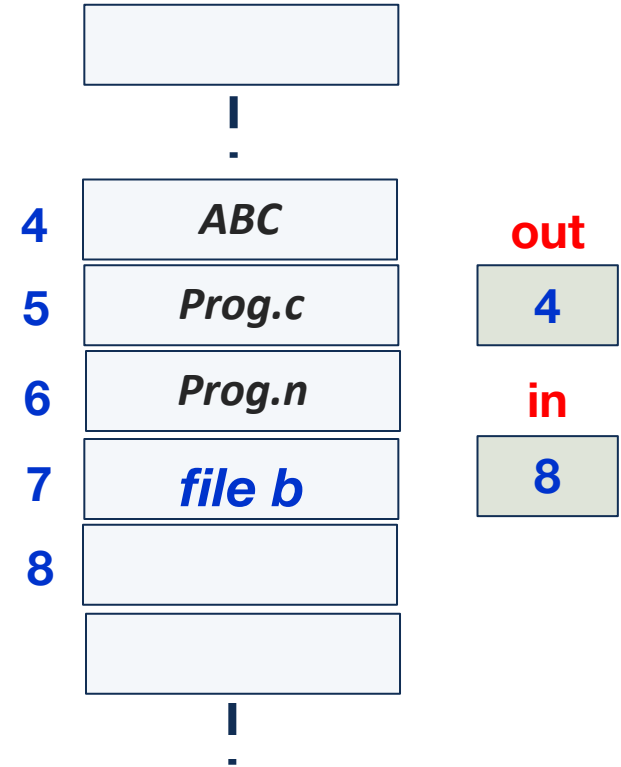
## Process A

read **in**  
store into A\_NextFreeSlot<sup>7</sup>  
add **file a** into the buffer  
**in** = A\_NextFreeSlot + 1

## Process B

read **in**  
store into B\_NextFreeSlot  
add **file b** into the buffer  
**in** = B\_NextFreeSlot + 1 ←

## Spooler Directory/ buffer



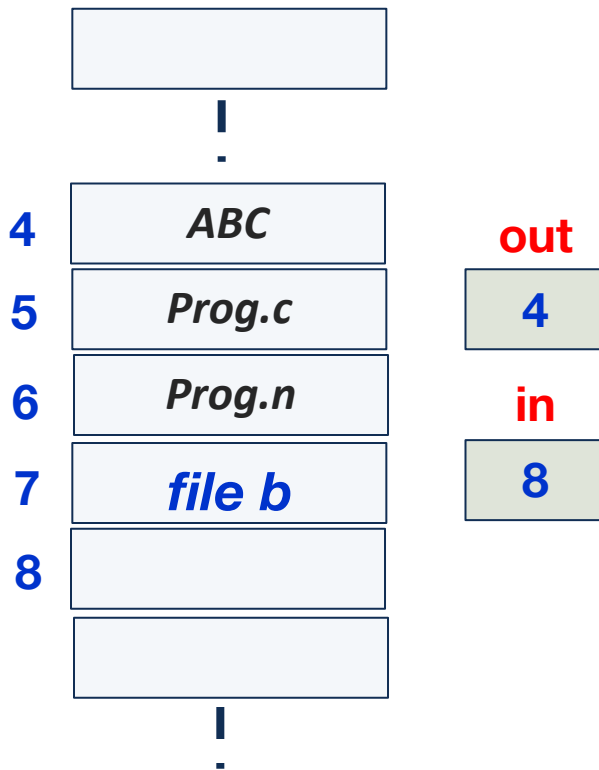
## Process A

read **in**  
store into A\_NextFreeSlot<sup>7</sup>  
add **file a** into the buffer  
**in** = A\_NextFreeSlot + 1

## Process B

read **in**  
store into B\_NextFreeSlot  
add **file b** into the buffer  
**in** = B\_NextFreeSlot + 1

## Spooler Directory/ buffer



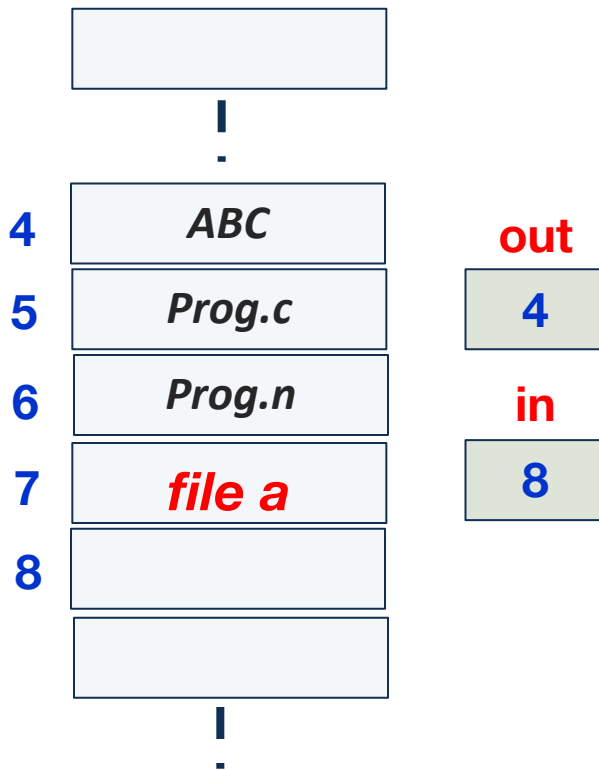
## Process A

read **in**  
store into A\_NextFreeSlot<sup>7</sup>  
add **file a** into the buffer ←  
**in** = A\_NextFreeSlot + 1

## Process B

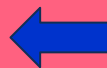
read **in**  
store into B\_NextFreeSlot  
add **file b** into the buffer  
**in** = B\_NextFreeSlot + 1

## Spooler Directory/ buffer



## Process A

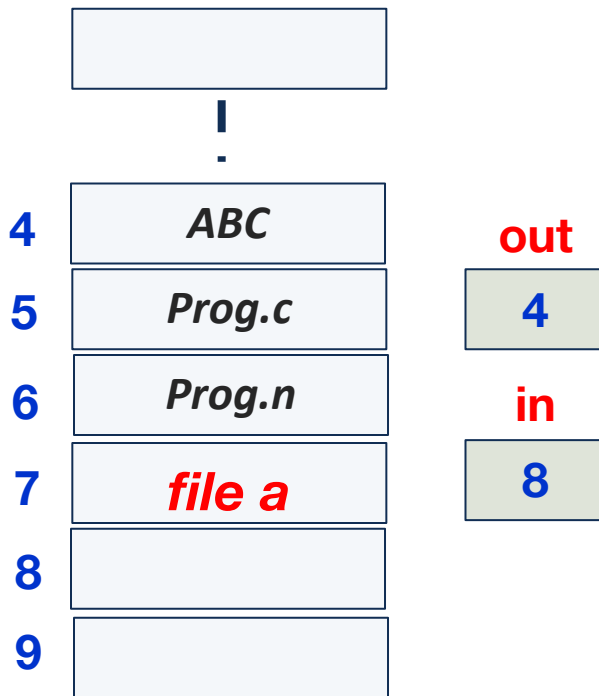
read in  
store into A\_NextFreeSlot  
add *file a* into the buffer  
in = A\_NextFreeSlot + 1



## Process B

read in  
store into B\_NextFreeSlot  
add *file b* into the buffer  
in = B\_NextFreeSlot + 1

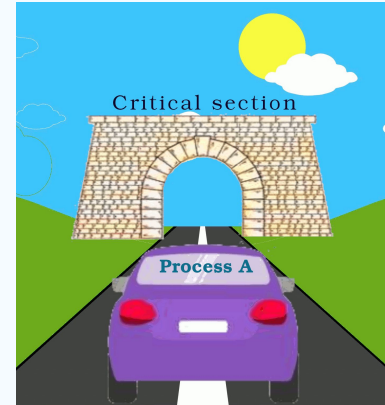
## Spooler Directory/ buffer



Process B will never  
receive any output

# The Critical Section Problem

- Each process has a segment of code called a **critical section**, in which the process may be changing common variables, updating a table, writing a file and so on.
- The important feature of the system is have to ensure that no other processes are executing its critical section when one particular process is executing its critical section.
- Processes cooperation protocol, whereby each process must request permission to enter its critical section. There is the entry section, then followed by exit section after that do the remaining code in the remainder section.



# The Critical Section Problem

Example:

Shared data:

Where is the Critical-section ?

```
int Balance;
```

Process A:

.....

enter section

Balance = Balance - 200;

exit section

remainder section

Process B:

.....

entry section

Balance = Balance - 100;

exit section

remainder section

# The Critical Section Problem

Example:

Shared data:

```
int Balance;
```

Process A:

.....

enter section

**Balance = Balance - 200;**

exit section

remainder section

Process B:

.....

entry section

**Balance = Balance - 100;**

exit section

remainder section



**Critical-section**



# Synchronization solutions

## Solution to Critical-Section requirements

1. **Mutual Exclusion** - If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections.
2. **Progress** - If no process is executing in its critical section and there exists some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely (unclear reasons).
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
  - Assume that each process executes at a nonzero speed
  - No assumption concerning relative speed of the  $N$  processes

# Peterson's Algorithm

## In Process synchronization

- Two process solution
- Assume that the LOAD and STORE instructions are atomic; that is, cannot be interrupted (privileged instructions).
- The two processes share two variables:
  - int **turn**;
  - Boolean **flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section.
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i]** = true implies that process  $P_i$  is ready!

Check the resources on C : [geeksforgeeks.org](https://www.geeksforgeeks.org)

# Critical Section

do {

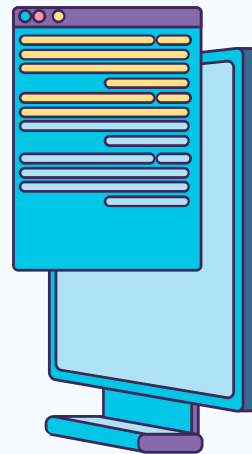
Entry section

Critical section

Exit section

Remainder section

} while (TRUE);



# Algorithm for Process $P_i$

do {

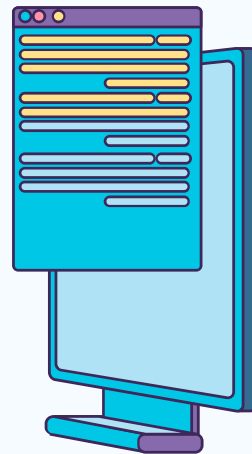
```
flag[i] = TRUE;           Entry section  
turn = j;  
while (flag[j] && turn == j);
```

Critical section

```
flag[i] = FALSE;         Exit section
```

Remainder section

} while (TRUE);



# Algorithm for Process $P_i$

**flag**

FALSE

FALSE

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] && turn == j);  
    Critical section  
    flag[i] = FALSE;  
    Remainder section  
} while (TRUE);
```

**$P_0$**

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] && turn == j);  
    Critical section  
    flag[i] = FALSE;  
    Remainder section  
} while (TRUE);
```

**$P_1$**

# Algorithm for Process $P_i$

**flag**

TRUE

FALSE

do {

flag[i] = TRUE;

turn = j;

while (flag[j] && turn == j);

Critical section

flag[i] = FALSE;

Remainder section

} while (TRUE);

**$P_0$**

do {

flag[i] = TRUE;

turn = j;

while (flag[j] && turn == j);

Critical section

flag[i] = FALSE;

Remainder section

} while (TRUE);

**$P_1$**

# Algorithm for Process $P_i$

**flag**

TRUE

FALSE

do {

    flag[i] = TRUE;

    turn = j;

 while (flag[j] && turn == j);

    Critical section

    flag[i] = FALSE;

    Remainder section

} while (TRUE);

**$P_0$**

do {

    flag[i] = TRUE;

    turn = j;

    while (flag[j] && turn == j);

    Critical section

    flag[i] = FALSE;

    Remainder section

} while (TRUE);

**$P_1$**

# Algorithm for Process $P_i$

**flag**

TRUE

TRUE

do {

flag[i] = TRUE;

turn = j;

 while (flag[j] && turn == j);

Critical section

flag[i] = FALSE;

Remainder section

} while (TRUE);

**$P_0$**

do {

flag[i] = TRUE;

turn = j;

 while (flag[j] && turn == j);

Critical section

flag[i] = FALSE;

Remainder section

} while (TRUE);

**$P_1$**



# Algorithm for Process $P_i$

**flag**

TRUE

TRUE

do {

flag[i] = TRUE;

turn = j;

 while (flag[j] && turn == j);

Critical section

flag[i] = FALSE;

Remainder section

} while (TRUE);

**$P_0$**

do {

flag[i] = TRUE;

turn = j;

 while (flag[j] && turn == j);

Critical section

flag[i] = FALSE;

Remainder section

} while (TRUE);

**$P_1$**

# Algorithm for Process $P_i$

**flag**

TRUE

TRUE

do {

    flag[i] = TRUE;

    turn = j;

    while (flag[j] && turn == j);

do {

    flag[i] = TRUE;

    turn = j;

    while (flag[j] && turn == j);

    Critical section

    flag[i] = FALSE;

    Remainder section

    Critical section

    flag[i] = FALSE;

    Remainder section

} while (TRUE);

} while (TRUE);

**$P_0$**

**$P_1$**

# Algorithm for Process $P_i$

**flag**

FALSE

TRUE

do {

    flag[i] = TRUE;

    turn = j;

    while (flag[j] && turn == j);

    Critical section

 flag[i] = FALSE;

    Remainder section

} while (TRUE);

**$P_0$**

do {

    flag[i] = TRUE;

    turn = j;

 while (flag[j] && turn == j);

    Critical section

    flag[i] = FALSE;

    Remainder section

} while (TRUE);

**$P_1$**

# Algorithm for Process $P_i$

**flag**

FALSE

TRUE

do {

    flag[i] = TRUE;

    turn = j;

    while (flag[j] && turn == j);

    Critical section

    flag[i] = FALSE;

    Remainder section

} while (TRUE);

**$P_0$**

do {

    flag[i] = TRUE;

    turn = j;

    while (flag[j] && turn == j);

    Critical section

    flag[i] = FALSE;

    Remainder section

} while (TRUE);

**$P_1$**

# Algorithm for Process $P_i$

**flag**

FALSE

TRUE

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] && turn == j);  
    Critical section  
    flag[i] = FALSE;  
    Remainder section  
} while (TRUE);
```

**$P_0$**

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] && turn == j);  
    Critical section  
    flag[i] = FALSE;  
    Remainder section  
} while (TRUE);
```

**$P_1$**

# Algorithm for Process $P_i$

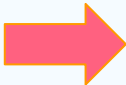
**flag**

FALSE

FALSE

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] && turn == j);  
    Critical section  
    flag[i] = FALSE;  
    Remainder section  
} while (TRUE);
```

**$P_0$**

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] && turn == j);  
    Critical section  
     flag[i] = FALSE;  
    Remainder section  
} while (TRUE);
```

**$P_1$**

# Algorithm for Process $P_i$

**flag**

FALSE

FALSE

do {

    flag[i] = TRUE;

    turn = j;

    while (flag[j] && turn == j);

    Critical section

    flag[i] = FALSE;

    Remainder section

} while (TRUE);

**$P_0$**

do {

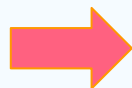
    flag[i] = TRUE;

    turn = j;

    while (flag[j] && turn == j);

    Critical section

    flag[i] = FALSE;



    Remainder section

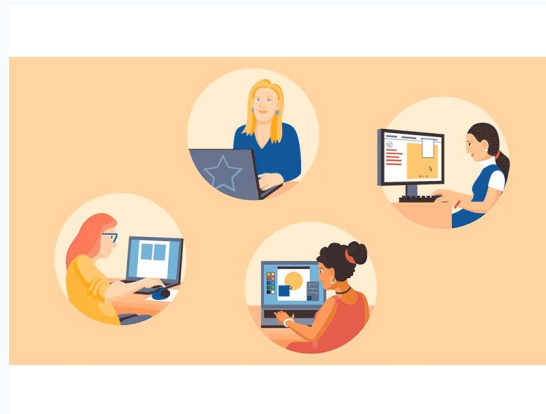
} while (TRUE);

**$P_1$**

# Synchronization Hardware (1/3)

## Disable Interrupt

- The critical-section problem could be solved simply in a uniprocessor environment, if we could prevent interrupts from occurring while a shared variable was being modified.
- The current sequence of instructions would be allowed to execute in order without preemption.
- This is often the approach taken by non preemptive kernel.

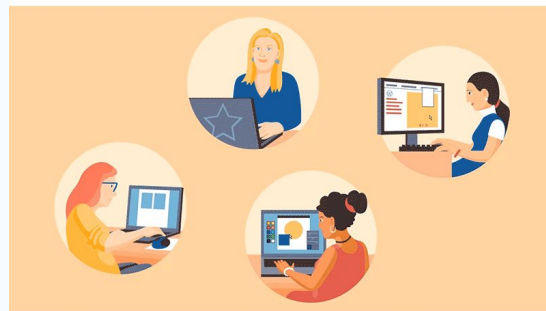




# Synchronization Hardware (2/3)

## Disable Interrupt

- a user-level program is given the ability to disable interrupts, then it can disable the timer interrupt and prevent context switching from taking place, thereby allowing it to use the processor without letting other processes execute.
- Unfortunately, this solution is not as feasible in a multiprocessors environment, because disabling interrupts on a multiprocessors can be time consuming as the message is passed to all the processors.

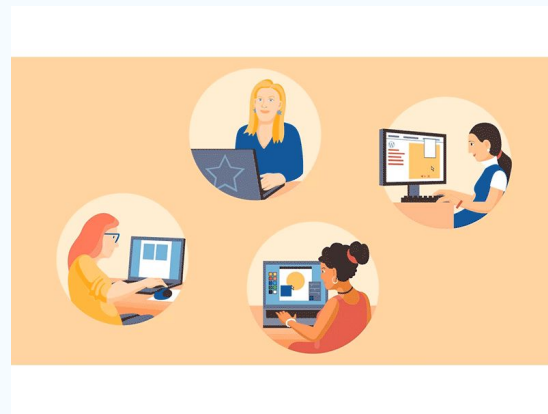


# Synchronization Hardware (3/3)

## Disable Interrupt

### why not disable interrupts?

- This message passing delays entry into each critical-section and the system efficiency decreases.
- Furthermore, disabling interrupts affect only the CPU that executed the disable instruction, where the other processor will continue running and can access the shared memory.



02

## Locks (Spin Locks)

A synchronization primitive that enforces limits on access to a resource when there are many threads of executions.



## Solution to Critical-section Problem Using Locks

Many systems provide hardware support for critical section code:

```
do {
```

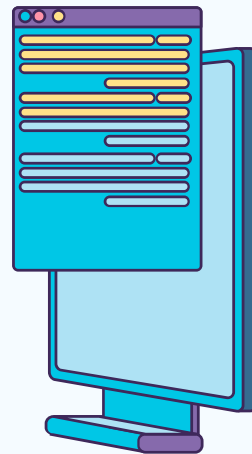
```
    acquire lock
```

```
    critical section
```

```
    release lock
```

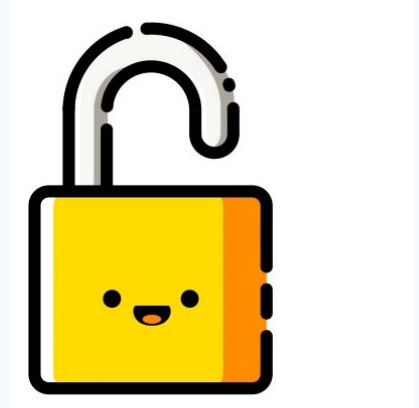
```
    remainder section
```

```
} while (TRUE);
```



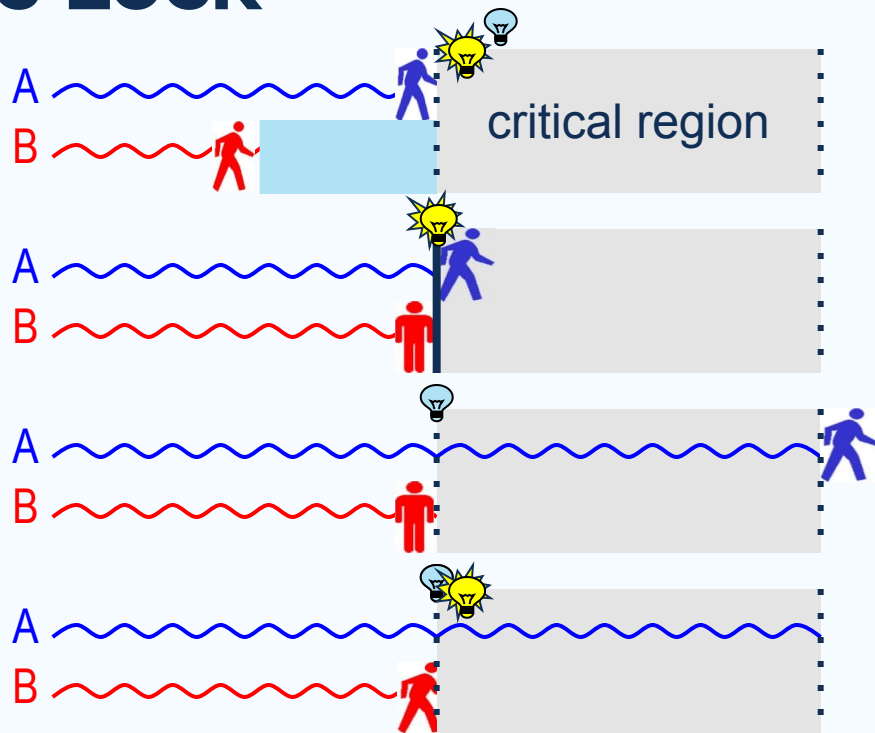
# Lock

- A single shared variable (lock) initially 0.
- When the process wants to enter its critical-section, it first tests the lock. If the lock is 0, the process set it to 1 and enter the critical-section.
- If the lock is already 1, the process just wait until it becomes 0.
- Disadvantages: Some threads/processes have to wait until a lock is released. If one of the threads holding a lock dies, stalls, blocks, or enters an infinite loop, other threads waiting for the lock may wait forever.



# Atomic Lock

1. thread A reaches CR and finds the lock at 0 and sets it in one shot, then enters
2. even if B comes right behind A, it will find that the lock is already at 1
3. thread A exits CR, then resets lock to 0
4. thread B finds the lock at 0 and sets it to 1 in one shot, just before entering CR



03

# Semaphore

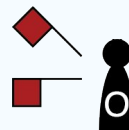
A variable or abstract data type used to control access to a common resource by multiple threads and avoid critical section problems.



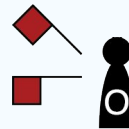
# Semaphore

- Semaphore  $S$  – integer variable
- Two standard operations modify  $S$ : `wait()` and `signal()`
  - Originally called `P()` and `V()`
- Less complicated
- Can only be accessed via **two indivisible (atomic) operations**

```
wait (S) {  
    while  $S \leq 0$ ;  
        // no-op  
     $S--$ ;  
}
```



```
signal (S) {  
     $S++$ ;  
}
```





# Semaphore as General Synchronization Tool

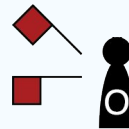
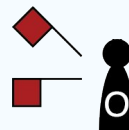
- Provide mutual exclusion

Semaphore S; //initialized to 1

wait (S);

**Critical Section**

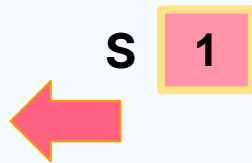
signal (S);



## Semaphore with two individual atomic operation

```
int balance, S;  
S = 1;  
// Process A  
...  
wait(S);  
balance = balance - 200;  
signal(S);  
remainder sessions
```

```
wait (S) {  
    while S ≤ 0 ; // no-op  
    S--;  
}
```



```
int balance, S;  
S = 1;  
// Process B  
...  
wait(S);  
balance = balance - 100;  
signal(S);  
remainder sessions
```

```
signal (S) {  
    S++;  
}
```

## Semaphore with two individual atomic operation

```
int balance, S;
```

```
S = 1;
```

```
// Process A
```

```
...
```

```
wait(S);
```

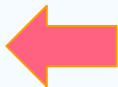
```
balance = balance - 200;
```

```
signal(S);
```

```
remainder sessions
```

S

0



```
int balance, S;
```

```
S = 1;
```

```
// Process B
```

```
...
```

```
wait(S);
```

```
balance = balance - 100;
```

```
signal(S);
```

```
remainder sessions
```

```
wait (S) {  
    while S ≤ 0 ; // no-op  
    S--;  
}
```

```
signal (S) {  
    S++;  
}
```

## Semaphore with two individual atomic operation

```
int balance, S;
```

```
S = 1;
```

```
// Process A
```

```
...
```

```
wait(S);
```

```
balance = balance - 200;
```

```
signal(S);
```

```
remainder sessions
```

**S**

**0**

```
int balance, S;
```

```
S = 1;
```

```
// Process B
```

```
...
```

```
wait(S);
```

```
balance = balance - 100;
```

```
signal(S);
```

```
remainder sessions
```

```
wait (S) {  
    while S ≤ 0 ; // no-op  
    S--;  
}
```

```
signal (S) {  
    S++;  
}
```

## Semaphore with two individual atomic operation

```
int balance, S;
```

```
S = 1;
```

```
// Process A
```

```
...
```

```
wait(S);
```

```
balance = balance - 200;
```

```
signal(S);
```

```
remainder sessions
```

**S**

**1**

```
int balance, S;
```

```
S = 1;
```

```
// Process B
```

```
...
```

```
wait(S);
```

```
balance = balance - 100;
```

```
signal(S);
```

```
remainder sessions
```

```
wait (S) {  
    while S ≤ 0 ; // no-op  
    S--;  
}
```

```
signal (S) {  
    S++;  
}
```

## Semaphore with two individual atomic operation

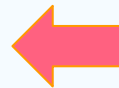
```
int balance, S;  
S = 1;  
// Process A  
...  
wait(S);  
balance = balance - 200;  
signal(S);  
remainder sessions
```

```
wait (S) {  
    while S ≤ 0 ; // no-op  
    S--;  
}
```

**S****1**

```
int balance, S;  
S = 1;  
// Process B  
...  
wait(S);  
balance = balance - 100;  
signal(S);  
remainder sessions
```

```
signal (S) {  
    S++;  
}
```



## Semaphore with two individual atomic operation

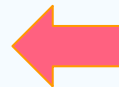
```
int balance, S;  
S = 1;  
// Process A  
...  
wait(S);  
balance = balance - 200;  
signal(S);  
remainder sessions
```

```
wait (S) {  
    while S ≤ 0 ; // no-op  
    S--;  
}
```

**S****0**

```
int balance, S;  
S = 1;  
// Process B  
...  
wait(S);  
balance = balance - 100;  
signal(S);  
remainder sessions
```

```
signal (S) {  
    S++;  
}
```



## Semaphore with two individual atomic operation

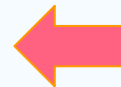
```
int balance, S;  
S = 1;  
// Process A  
...  
wait(S);  
balance = balance - 200;  
signal(S);  
remainder sessions
```

```
wait (S) {  
    while S ≤ 0 ; // no-op  
    S--;  
}
```

**S****0**

```
int balance, S;  
S = 1;  
// Process B  
...  
wait(S);  
balance = balance - 100;  
signal(S);  
remainder sessions
```

```
signal (S) {  
    S++;  
}
```





## Semaphore with two individual atomic operation

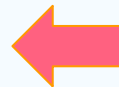
```
int balance, S;  
S = 1;  
// Process A  
...  
wait(S);  
balance = balance - 200;  
signal(S);  
remainder sessions
```

```
wait (S) {  
    while S ≤ 0 ; // no-op  
    S--;  
}
```

**S****1**

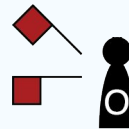
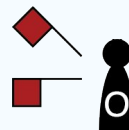
```
int balance, S;  
S = 1;  
// Process B  
...  
wait(S);  
balance = balance - 100;  
signal(S);  
remainder sessions
```

```
signal (S) {  
    S++;  
}
```



# Semaphore Implementation

- Must guarantee that no two processes can execute `wait ()` and `signal ()` on the same semaphore at the same time
- Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section.
- Could now have busy waiting in critical section implementation
  - But implementation code is short
  - **Little busy waiting** if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution.



# Semaphore Bounded-Buffer

$N$  buffers, each can hold one item

## Semaphore Mutex

Initialise to the  
value 1

## Semaphore Full

Initialise to the  
value 0

## Semaphore Empty

Initialise to the  
value  $N$

# Semaphore in Producer-Consumer

## Structure Process of Producer

```
do {  
    // produce an item  
    wait (empty);  
    wait (mutex);  
  
    // add item to buffer  
  
    signal (mutex);  
    signal (full);  
} while (true);
```

## Structure Process of Consumer

```
do {  
    wait (full);  
    wait (mutex);  
    // remove item from buffer  
  
    signal (mutex);  
    signal (empty);  
  
    // consume the removed item  
} while (true);
```

**mutex**

1

**empty**

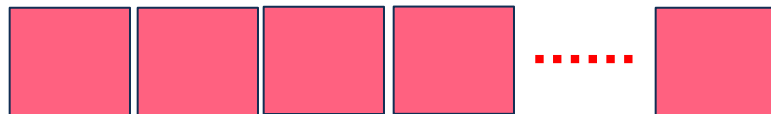
n

**full**

0

Next  
P

Buffer



0

1

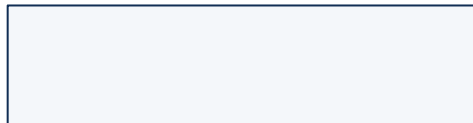
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    Wait(empty);  
    Wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    Wait(full);  
    Wait(mutex);  
    //remove item from buffer to nextc  
    Signal(mutex);  
    Signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

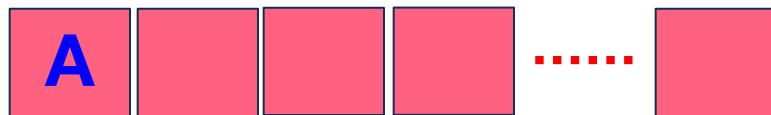
n

**full**

0

Next  
P

**Buffer**



0

1

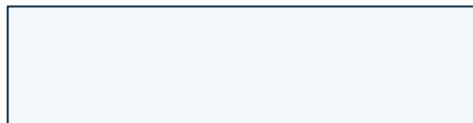
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    Wait(empty);  
    Wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

n-1

**full**

0

Next  
P

**Buffer**



0

1

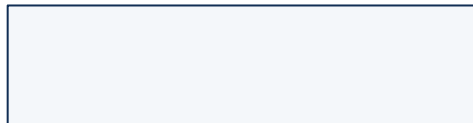
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    Wait(empty);  
    Wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

0

**empty**

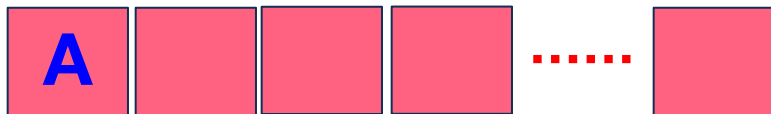
n-1

**full**

0

Next  
P

**Buffer**



0

1

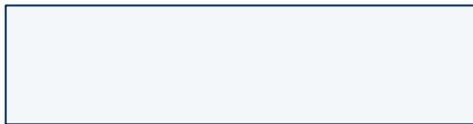
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    Wait(empty);  
    Wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```



**mutex**

0

**empty**

n-1

**full**

0

Next  
P



**Buffer**



0

1

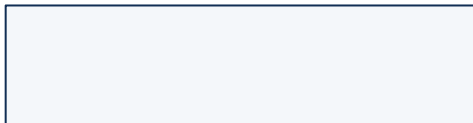
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    Wait(empty);  
    Wait(mutex);  
    ➔ //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

n-1

**full**

0

Next  
P



**Buffer**



0

1

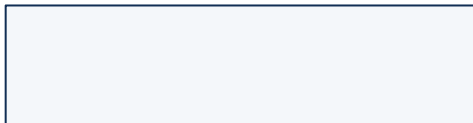
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    Wait(empty);  
    Wait(mutex);  
    //add nextp to buffer  
    → signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

n-1

**full**

1

Next  
P



**Buffer**



0

1

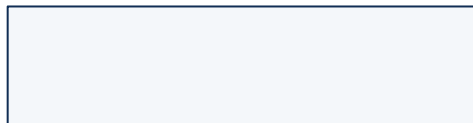
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

n-1

**full**

1

Next  
P



**Buffer**



0

1

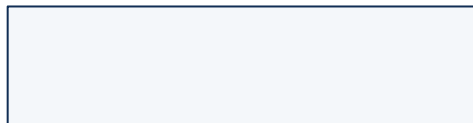
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

n-2

**full**

1

Next  
P



**Buffer**



0

1

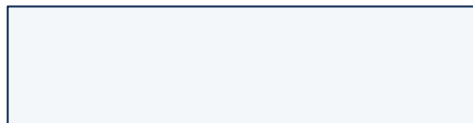
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

0

**empty**

n-2

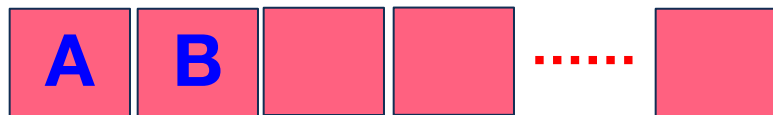
**full**

1

**Next  
P**



**Buffer**



0

1

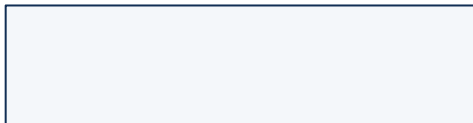
2

3

n

**Next  
C**

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

0

**empty**

n-2

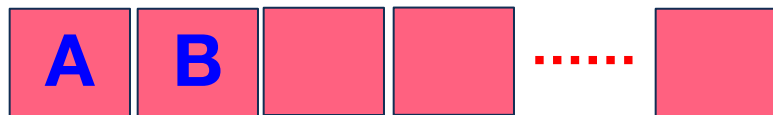
**full**

1

Next  
P



**Buffer**



0

1

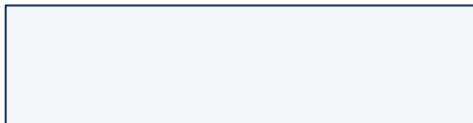
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    ➔ //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

n-2

**full**

1

Next  
P



**Buffer**



0

1

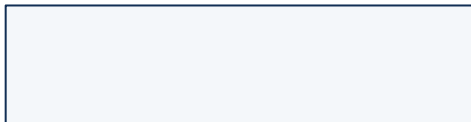
2

3

n

Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    → signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```



**mutex**

1

**empty**

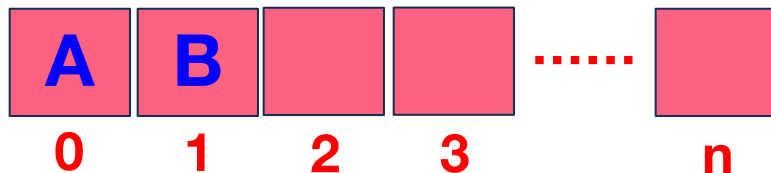
n-2

**full**

2

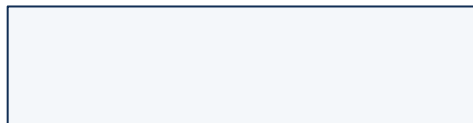
Next  
P

**Buffer**



Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

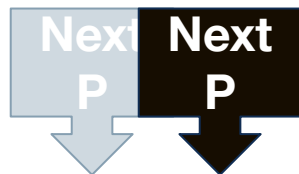
1

**empty**

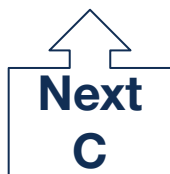
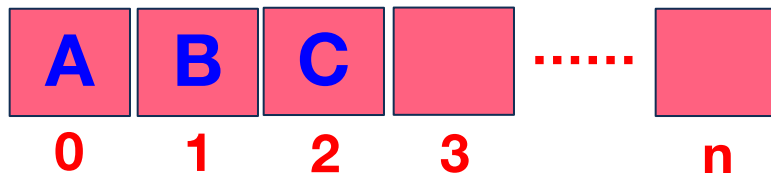
n-3

**full**

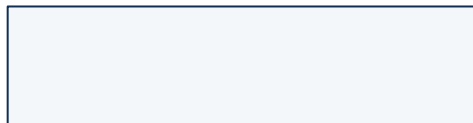
3



**Buffer**



**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    wait(full);  
    wait(mutex);  
    //remove item from buffer to nextc  
    signal(mutex);  
    signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

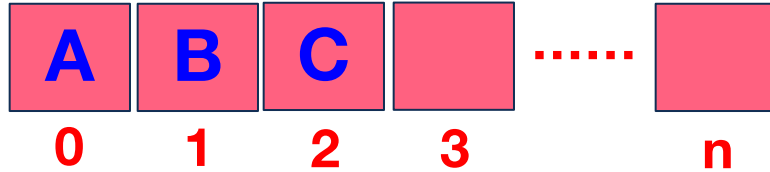
$n-3$

**full**

2

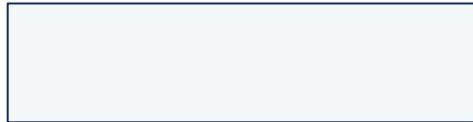
Next  
P

**Buffer**



Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    Wait(full);  
    Wait(mutex);  
    //remove item from buffer to nextc  
    Signal(mutex);  
    Signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

0

**empty**

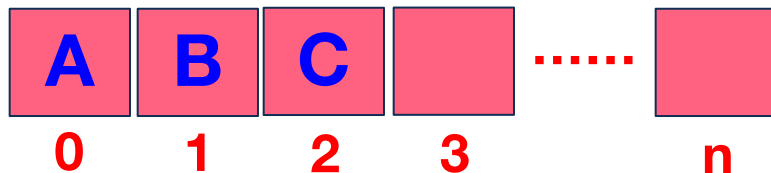
n-3

**full**

2

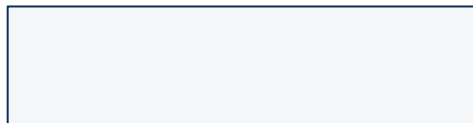
Next  
P

**Buffer**



Next  
C

**Display**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    Wait(full);  
    Wait(mutex);  
    //remove item from buffer to nextc  
    Signal(mutex);  
    Signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

0

**empty**

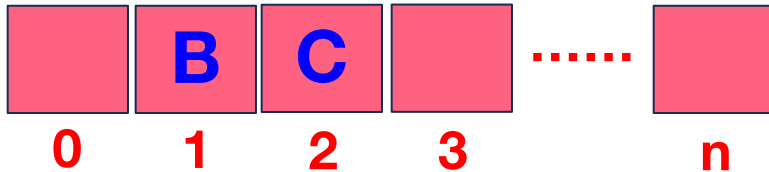
n-3

**full**

2

Next  
P

**Buffer**



Next  
C

**Display**

A

**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    Wait(full);  
    Wait(mutex);  
    //remove item from buffer to nextc  
    Signal(mutex);  
    Signal(empty);  
    //consume the item in nextc  
}(true)
```

# mutex

1

# empty

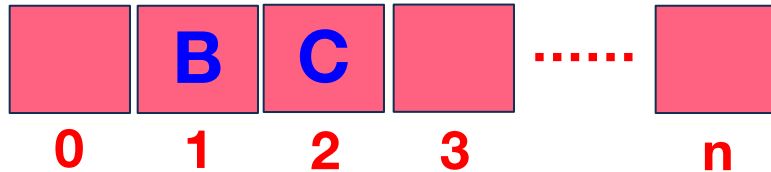
n-3

# full

2

Next  
P

Buffer



Next  
C

Display

A

## Producer

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

## Consumer

```
do{  
    Wait(full);  
    Wait(mutex);  
    //remove item from buffer to nextc  
    Signal(mutex);  
    Signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

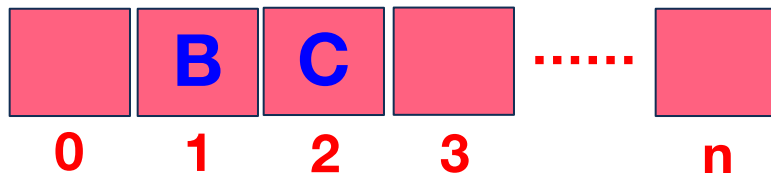
n-2

**full**

2

Next  
P

**Buffer**



Next  
C

**Display**

A

**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    Wait(full);  
    Wait(mutex);  
    //remove item from buffer to nextc  
    Signal(mutex);  
    Signal(empty);  
    //consume the item in nextc  
}(true)
```

**mutex**

1

**empty**

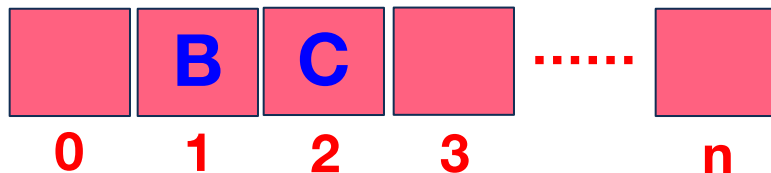
n-2

**full**

2

Next  
P

**Buffer**



Next  
C

**Display**

A

**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    Wait(full);  
    Wait(mutex);  
    //remove item from buffer to nextc  
    Signal(mutex);  
    Signal(empty);  
    //consume the item in nextc  
}(true)
```



**mutex**

1

**empty**

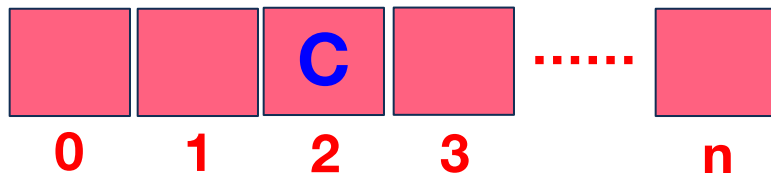
n-1

**full**

1

Next  
P

**Buffer**



**Producer**

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

**Consumer**

```
do{  
    Wait(full);  
    Wait(mutex);  
    //remove item from buffer to nextc  
    Signal(mutex);  
    Signal(empty);  
    //consume the item in nextc  
}(true)
```

# mutex

1

# empty

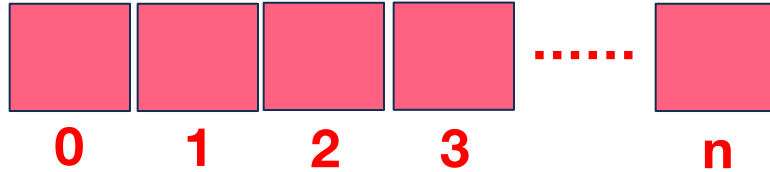
n

# full

0

Next  
P

Buffer



Next  
C

Next  
C

Display

A B C

## Producer

```
do{  
    //produce an item in nextp  
    wait(empty);  
    wait(mutex);  
    //add nextp to buffer  
    signal(Mutex);  
    signal(full);  
}(true)
```

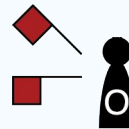
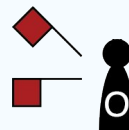
## Consumer

```
do{  
    Wait(full);  
    Wait(mutex);  
    //remove item from buffer to nextc  
    Signal(mutex);  
    Signal(empty);  
    //consume the item in nextc  
}(true)
```

# Semaphore Implementation

## With No Busy Waiting

- With each semaphore there is an associated waiting queue. Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list
- Two operations:
  - **Block()** – place the process invoking the operation on the appropriate waiting queue.
  - **Wakeup()** – remove one of processes in the waiting queue and place it in the ready queue.



# Semaphore implementation

## Implementation of Wait()

```
wait (S){  
    value--;  
  
    if (value < 0) {  
        block();  
  
        -- add this process to  
           waiting queue  
    }
```

## Implementation of Signal()

```
Signal (S){  
    value++;  
  
    if (value ≤ 0) {  
        wakeup(P);    }  
  
    -- remove a process P from  
       the waiting queue  
}
```

ooo

Mutex  
value

1

Mutex list

### Process B

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

### Process C

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

### Process A

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

```
wait (S){  
    value--;  
    if (value < 0) {  
        //add this to waiting  
        queue  
        block();  
    }  
}
```

**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

1

Mutex list

**Process B**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Process C**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Process A**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

```
wait (S){  
    value--;  
    if (value < 0) {  
        //add this to waiting  
        queue  
        block();  
    }  
}
```

**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

0

Mutex list

**Process B**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Process C**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Process A**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

```
wait (S){  
    value--;  
    if (value < 0) {  
        //add this to waiting  
        queue  
        block();  
    }  
}
```

**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

0

Mutex list

**Process B**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Process C**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Process A**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder



```
wait (S){  
    value--;  
    if (value < 0) {  
        //add this to waiting  
        queue  
        block();  
    }  
}
```

**Semaphore implementation with No Busy Waiting**



ooo

Mutex  
value

0

Mutex list

### Process B

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

### Process C

Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

```
wait (S){  
    value--;  
    if (value < 0) {  
        //add this to waiting  
        queue  
        block();  
    }  
}
```

**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

-1

Mutex list

### Process B

Start

Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

### Process C

Start

Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

Process A

Start

Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

```
wait (S){  
    value--;  
    if (value < 0) {  
        //add this to waiting  
        queue  
        block();  
    }  
}
```

**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

-1

Mutex list

Process B

Blocked

Process B

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

Process C

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

```
wait (S){  
    value--;  
    if (value < 0) {  
        //add this to waiting  
        queue  
        block();  
    }  
}
```

Semaphore implementation with No Busy Waiting

Mutex  
value

**-1**

Mutex list

**Process B**

**Blocked**

**Process B**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder



**Process C**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

```
wait (S){
    value--;
    if (value < 0) {
        //add this to waiting
        queue
        block();
    }
}
```

**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

-2

Mutex list

Process B

**Blocked**

**Process B**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder



**Process C**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

```
wait (S){  
    value--;  
    if (value < 0) {  
        //add this to waiting  
        queue  
        block();  
    }  
}
```



**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

**-2**

Mutex list

**Process B, Process C**

**Blocked**

**Process B**

Start

Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

**Blocked**

**Process C**

Start

Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

```
wait (S){  
    value--;  
    if (value < 0) {  
        //add this to waiting  
        queue  
        block();  
    }  
}
```

**Semaphore implementation with No Busy Waiting**

Mutex  
value

**-1**

Mutex list

**Process B, Process C**

**Blocked**

**Process B**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Process A**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Blocked**

**Process C**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

```
signal (S){
    value++;
    if (value <= 0) {
        //remove a process P
        from the waiting queue
        wakeup(P);
    }
}
```

**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

-1

Mutex list

**Process B, Process C**

**Blocked**

**Process B**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Blocked**

**Process C**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

**Process A**

Start

Wait(Mutex);

Critical Section

Signal(Mutex);

Do remainder

```
signal (S){  
    value++;  
    if (value <= 0) {  
        //remove a process P  
        from the waiting queue  
        wakeup(P);  
    }  
}
```

**Semaphore implementation with No Busy Waiting**



ooo

Mutex  
value

0

Mutex list

Process C

### Process B

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

Blocked

### Process C

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

```
signal (S){  
  value++;  
  if (value <= 0) {  
    //remove a process P  
    from the waiting queue  
    wakeup(P);  
  }  
}
```

**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

0

Mutex list

Process C

### Process B

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

Blocked

### Process C

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

```
signal (S){  
    value++;  
    if (value <= 0) {  
        //remove a process P  
        from the waiting queue  
        wakeup(P);  
    }  
}
```

**Semaphore implementation with No Busy Waiting**

Mutex  
value

1

Mutex list

### Process B

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

### Process C

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

```
signal (S){
  value++;
  if (value <= 0) {
    //remove a process P
    from the waiting queue
    wakeup(P);
  }
}
```

**Semaphore implementation with No Busy Waiting**

ooo

Mutex  
value

1

Mutex list

### Process B

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder

### Process C

Start  
Wait(Mutex);  
Critical Section  
Signal(Mutex);  
Do remainder



```
signal (S){  
  value++;  
  if (value <= 0) {  
    //remove a process P  
    from the waiting queue  
    wakeup(P);  
  }  
}
```



**Semaphore implementation with No Busy Waiting**

# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

```
wait (S){
    value--;

    if (value < 0) {
        //add this process to
        waiting queue

        block();
    }
}
```

	S	Q
value	1	1

$P_0$   
wait (S);  
wait (Q);

·  
·  
signal (S);  
signal (Q);

$P_1$   
wait (Q);  
wait (S);

·  
·  
signal (Q);  
signal (S);

# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

	<b>S</b>	<b>Q</b>
<b>value</b>	0	1

```
wait (S){
    value--;

    if (value < 0) {
        //add this process to
        waiting queue

        block();
    }
}
```

$P_0$	$P_1$
wait (S);	wait (Q);
wait (Q);	wait (S);
.	.
.	.
signal (S);	signal (Q);
signal (Q);	signal (S);

# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

```
wait (S){
    value--;

    if (value < 0) {
        //add this process to
        waiting queue

        block();
    }
}
```

	S	Q
value	0	0

$P_0$	$\rightarrow$	$P_1$
wait (S);		wait (Q);
wait (Q);		wait (S);
.		.
.		.
signal (S);		signal (Q);
signal (Q);		signal (S);

# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

	<b>S</b>	<b>Q</b>
<b>value</b>	0	-1

```
wait (S){
    value--;

    if (value < 0) {
        //add this process to
        waiting queue

        block();
    }
}
```

$P_0$	$P_1$
wait (S);	wait (Q);
wait (Q);	wait (S);
.	.
.	.
signal (S);	signal (Q);
signal (Q);	signal (S);



# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

```
wait (S){
    value--;

    if (value < 0) {
        //add this process to
        waiting queue

        block();
    }
}
```

	S	Q
value	-1	-1

$P_0$	$P_1$
wait (S);	wait (Q);
wait (Q);	wait (S);
.	.
.	.
signal (S);	signal (Q);
signal (Q);	signal (S);

# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

```

signal (S){
    value++;

    if (value <= 0) {
        //remove a process P
        from the waiting queue

        wakeup(P);
    }
}

```

	<b>S</b>	<b>Q</b>
<b>value</b>	-1	-1

$P_0$   
wait (S);  
wait (Q);

$P_1$   
wait (Q);  
wait (S);

.

.

.

.



→ signal (S);  
signal (Q);

signal (Q);  
signal (S);

# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

	<b>S</b>	<b>Q</b>
<b>value</b>	1	1

```
wait (S) {
    while S ≤ 0 ; // no-op
    S--;
}
```

$P_0$   
wait (S);  
wait (Q);

·  
·  
signal (S);  
signal (Q);

$P_1$   
wait (Q);  
wait (S);

·  
·  
signal (Q);  
signal (S);

# Deadlock and Starvation

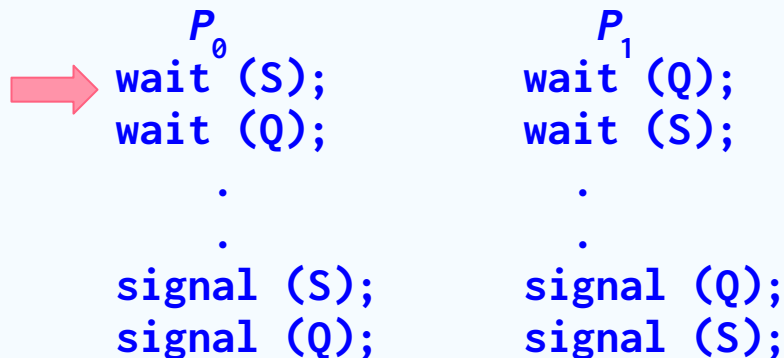
The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

	<b>S</b>	<b>Q</b>
<b>value</b>	0	1

```
wait (S) {
    while S ≤ 0 ; // no-op
    S--;
}
```



# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

```
wait (S) {
    while S ≤ 0 ; // no-op
    S--;
}
```

	S	Q
value	0	0

$P_0$   
wait (S);  
wait (Q);

·  
·  
signal (S);  
signal (Q);

→  $P_1$   
wait (Q);  
wait (S);

·  
·  
signal (Q);  
signal (S);

# Deadlock and Starvation

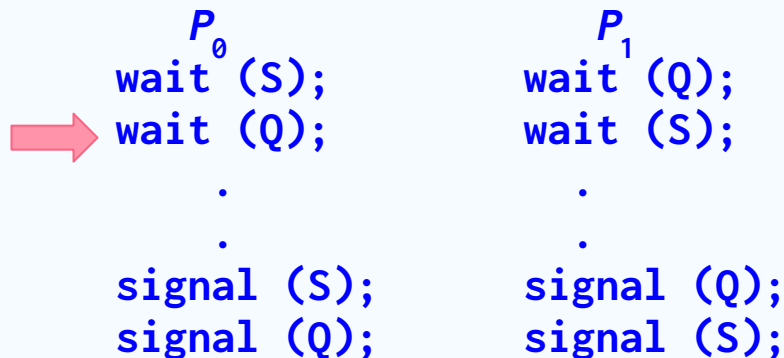
The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

	<b>S</b>	<b>Q</b>
<b>value</b>	0	0

```
wait (S) {  
    while  $S \leq 0$  ; // no-op  
    S--;  
}
```



# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

	<b>S</b>	<b>Q</b>
<b>value</b>	0	0

```
wait (S) {
    while S ≤ 0 ; // no-op
    S--;
}
```

$P_0$   
wait (S);  
wait (Q);

·  
·  
signal (S);  
signal (Q);

$P_1$   
wait (Q);  
wait (S);

→

·  
·  
signal (Q);  
signal (S);

# Deadlock and Starvation

The implementation of a semaphore with a **waiting queue** may result in a situation where two or more processes are waiting for each other.

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let **S** and **Q** be two semaphores initialized to 1

	<b>S</b>	<b>Q</b>
<b>value</b>	0	0

```
signal (S) {
    S++;
}
```

$P_0$   
wait (S);  
wait (Q);

$P_1$   
wait (Q);  
wait (S);



signal (S);  
signal (Q);

signal (Q);  
signal (S);



# Explanation



Suppose that P0 executes wait(S) and then P1 executes wait(Q).

When P0 executes wait(Q), it must wait until P1 executes signal(Q).



Similarly, when P1 executes wait(S), it must wait P0 executes signal(S).

Since these signal() operations cannot be executed, P0 and P1 are deadlock.

# Problems with semaphore

**Incorrect use of semaphore operations:**

**01** **signal (mutex) .... wait (mutex)**

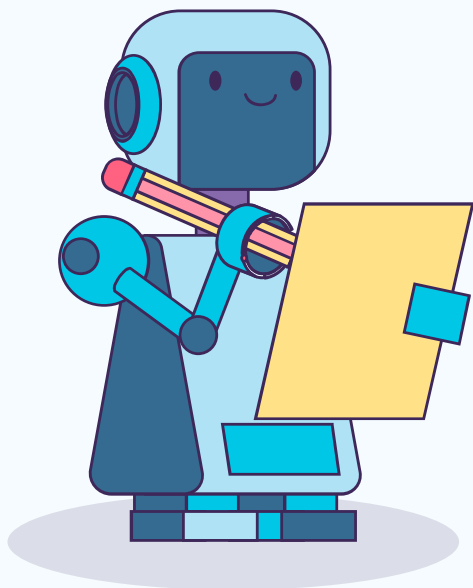
Several processes may execute in the CS

**02** **wait (mutex) ... wait (mutex)**

Deadlock will occur

**03** **Omitting of wait (mutex) or/and signal (mutex)**

Several processes may execute in the CS



# Thank you

CREDITS: This presentation template was created by **Slidesgo**, and includes icons by **Flaticon**, and infographics & images by **Freepik**